

Systems Biology Graphical Notation: Process Description language Level 1

Version 2.0

Date: March 16, 2013

Disclaimer: This is a working draft of the SBGN Process Description Level 1 Version 2.0 specification. It is not a normative document.

Editors:

Stuart Moodie	<i>EMBL European Bioinformatics Institute, UK</i>
Nicolas Le Novère	<i>EMBL European Bioinformatics Institute, UK</i>
Emek Demir	<i>Sloan-Kettering Institute, USA</i>
Huaiyu Mi	<i>University of Southern California, USA</i>
Alice Villéger	<i>Freelance IT Consultant, London, UK</i>

To discuss any aspect of SBGN, please send your messages to the mailing list sbgn-discuss@sbgn.org. To get subscribed to the mailing list or to contact us directly, please write to sbgn-editors@lists.sourceforge.net. Bug reports and specific comments about the specification should be entered in the issue tracker http://p.sf.net/sbgn/pd_tracker.



Contents

Contents	ii
I SBGN Primer	1
1 Introduction	2
1.1 SBGN levels and versions	3
1.2 Developments, discussions, and notifications of updates	4
1.3 How to read this document	4
2 Concepts	6
2.1 Overview of SBGN Process Descriptions	6
2.2 Some more concept text	9
3 Language Overview	12
3.1 Symbols used in SBGN Process Descriptions	12
3.2 Building a SBGN Process Description map	39
4 Worked Examples	44
II Technical Reference	45
5 Language Elements	46
5.1 Introduction	46
5.2 Note on typographical convention	47
5.3 How to read the Language Specification	47
5.4 Overview	50
5.5 Definitions	52
5.6 Controlled vocabularies	119

5.7	Entity Pool Node Identity and Cloning	121
5.8	Map and Submap Linking	122
6	Layout Rules for a Process Description	127
6.1	Introduction	127
6.2	Requirements	128
6.3	Recommendations	130
6.4	Additional suggestions	131
7	Acknowledgments	133
7.1	Level 1 Release 1.0	133
7.2	Level 1 Release 1.1	133
7.3	Level 1 Release 1.2	133
7.4	Level 1 Release 1.3	134
7.5	Comprehensive list of acknowledgements	134
7.6	Financial Support	134
A	Complete examples of Process Description Maps	135
B	Reference card	141
C	Validation and consistency rules	143
C.1	EPN rules	143
D	Issues postponed to future levels	144
D.1	Multicompartment entities	144
D.2	Logical combination of state variable values	145
D.3	Non-chemical entity nodes	145
D.4	Generics	145
D.5	State and transformation of compartments	145
E	Revision History	146
E.1	Version 1.0 to Version 1.1	146
E.2	Version 1.1 to Version 1.2	147
E.3	Version 1.2 to Version 1.3	148

Part I

SBGN Primer

Introduction

With the rise of systems and synthetic biology, the use of graphical representations of pathways and networks to describe biological systems has become pervasive. It was therefore important to use a consistent notation that would allow people to interpret those maps easily and quickly, without the need of extensive legends. Furthermore, distributed investigation of biological systems in different labs as well as activities like synthetic biology, that reconstruct biological systems, need to exchange their descriptions unambiguously, as engineers exchange circuit diagrams.

The goal of the Systems Biology Graphical Notation (SBGN) is to standardize the graphical/visual representation of biochemical and cellular processes. SBGN defines comprehensive sets of symbols with precise semantics, together with detailed syntactic rules defining their use. It also describes the manner in which such graphical information should be interpreted. SBGN is made up of three different and complementary languages [?]. This document defines the *Process Description* visual language of SBGN. Process Descriptions are one of three views of a biological process offered by SBGN. It is the product of many hours of discussion and development by many individuals and groups.

What are the languages?

PD is a language that permits the description of all the processes taking place in a biological system. The ensemble of all these processes constitute a Description. **ER** is a language that permits the description of all the relations involving the entities of a biological system. The ensemble of all these relations constitute a Relationship. **AF** is a language that permits the description of the flow of activity in a biological system.

Nomenclature

The three languages of SBGN should be referred to as:

- the Process Description language.
- the Entity Relationship language.
- the Activity Flow language.

Abbreviated as:

- the PD language. 5
- the ER language.
- the AF language.

A specific representation of a biological system in one of the SBGN languages should be referred to as:

- a Process Description map. 10
- an Entity Relationship map.
- an Activity Flow map.

Abbreviated as:

- a PD map.
- an ER map. 15
- an AF map.

The corpus of all SBGN representations should be referred to as:

- Process Descriptions.
- Entity Relationships.
- Activity Flows. 20

The capitalization is important. PD, ER and AF are names of languages. As such they must be capitalized in English. This is not the case of the accompanying noun (language or map).

1.1 SBGN levels and versions

It was clear at the outset of SBGN development that it would be impossible to design a perfect and complete notation right from the beginning. Apart from the prescience this would require (which, sadly, none of the authors possess), it also would likely need a vast language that most newcomers would shun as being too complex. Thus, the SBGN community followed an idea used in the development of other standards, i.e. stratify language development into levels. 25 30

A *level* of one of the SBGN languages represents a set of features deemed to fit together cohesively, constituting a usable set of functionality that the user community agrees is sufficient for a reasonable set of tasks and goals. Within *levels*, *versions* represent evolutions of a language, that may involve new glyphs, refined semantics, but no fundamental change of the way maps are to be generated and interpreted. In addition new versions should be backwards compatible, i.e., Process Description maps that conform to an earlier version of the Process Description language within the same level should still be valid. This does not apply to a new levels. 35

Capabilities and features that cannot be agreed upon and are judged insufficiently critical to require inclusion in a given level, are postponed to a higher 40

level or version. In this way, the development of SBGN languages is envisioned to proceed in stages, with each higher levels adding richness compared to the levels below it.

1.2 Developments, discussions, and notifications of updates

5

The SBGN website (<http://sbgn.org/>) is a portal for all things related to SBGN. It provides a web forum interface to the SBGN discussion list (sbgn-discuss@caltech.edu) and information about how anyone may subscribe to it. The easiest and best way to get involved in SBGN discussions is to join the mailing list and participate.

Face-to-face meetings of the SBGN community are announced on the website as well as the mailing list. Although no set schedule currently exists for workshops and other meetings, we envision holding at least one public workshop per year. As with other similar efforts, the workshops are likely to be held as satellite workshops of larger conferences, enabling attendees to use their international travel time and money more efficiently.

10

15

Notifications of updates to the SBGN specification are also broadcast on the mailing list and announced on the SBGN website.

1.3 How to read this document

SBGN is a visual language aimed at biologists and those working in the life sciences in general. While you could use pen and paper to draw diagram in Process Description we expect that most of you will use specialist software tools most of the time. Often software tools will be developed by software engineers who do not have a deep understanding of biology and biological systems. Therefore, the challenge for us in drafting this specification is to satisfy both readerships, biologists and software engineers. Our answer to this dilemma has been to provide a number of examples that we explain in detail. By working through the examples the fine detail of Process Description will become clear to all. Finally we provide a technical reference at the end of this document that provides a detailed and unambiguous specification of the language. To help you get started, we have come up with the following advice on how each readership could read this specification.

20

25

30

You may be a user, for example a life-scientist who is interested mainly in using Process Description language. In addition, you may want to understand what it can do for and also suggest ways that the language can be improved. In either case we recommend that you start with the language overview in chapter ?? on page ?. From there you should move onto the examples in chapter ?? (page ?). Once you are confident with the material in the examples, you may want to dip in and out of the technical reference (Part ?? on page ?), but this is not essential.

35

If you are a software engineer then we recommend that you start with the language overview in chapter ?? on page ?. After this, work through the examples in

chapter ?? on page ?? and try to understand the language features in this way. Finally, when it comes to implementing Process Description language support you will find the technical reference (Part ?? on page ??) very important, but only after you have an understanding of Process Description from this primer.

Concepts

2.1 Overview of SBGN Process Descriptions

To briefly describe what SBGN Process Description language is about, let's give a brief overview of some of the relevant concepts with the help of an example shown in Figure 2.1. It is a simple map for part of a mitogen-activated protein kinase (MAPK) cascade. The larger nodes in the figure (some of which are in the shape of rounded rectangles and others in the shape of circles) represent biological materials—things like macromolecules and simple chemicals (NB: the nodes representing physical entities (or proxies to physical entities) will always be colored in yellow in this document. Color is not part of the SBGN specification though). The biological materials are altered via processes (colored in green in this document), which are indicated in Process Description language by lines with arrows and other decorations. In this particular map, all of the processes happen to be the same: processes catalyzed by biochemical entities. The directions of the arrows indicate the direction of the processes; for example, unphosphorylated RAF kinase proceeds to phosphorylated RAF kinase via a process catalyzed by RAS. Although ATP and ADP are shown as incidental to the phosphorylations on this particular graph, they are involved in the same process as the proteins getting phosphorylated. The small circles on the nodes for RAF and other entity pools represent state variables (in this case, phosphorylation sites).

The essence of the Process Descriptions is *change*: it shows how different entities in the system process from one form to another. The entities themselves can be many different things. In the example of Figure 2.1, they are either pools of macromolecules or pools of simple chemicals, but as will become clear later in this chapter, they can be other conceptual and material constructs as well. Note also that we speak of *entity pools* rather than individuals; this is because in biochemical network models, one does not focus on single molecules, but rather collections of molecules of the same kind. The molecules in a given pool are considered indistinguishable from each other. The way in which one type of entity is transformed into another is conveyed by a *process node* and arcs between entity pool




Figure 2.1: This example of a Process Description uses two kinds of entity pool nodes: one for pools of different macromolecules (Section 5.5) and another for pools of simple chemicals (Section 5.5). Most macromolecule nodes in this map are adorned with state variables (Section 5.5) representing phosphorylation states. This map uses one type of process node, the process node (Section 5.5), and three kind of connecting arc, consumption (Section 5.5), production (Section 5.5) and catalysis (Section 5.5). Finally, some entity pool nodes have dark bands along their bottoms; these are clone markers (Section 5.5) indicating that the same pool nodes appear multiple times in the map.

nodes and process nodes indicate an influence by the entities on the processes. In the case of Figure 2.1, those arcs describe consumption (Section 5.5), production (Section 5.5) and catalysis (Section 5.5), but others are possible. Finally, nodes in Process Descriptions are usually not repeated; if they do need to be repeated, they are marked with *clone markers*—specific modifications to the appearance of the node (Section 5.5). The details of this and other aspects of Process Description notation are explained in the rest of this chapter. 5

A reference card depicting all the symbols of SBGN Process Description Level 1 is present at the end of this document.

Lets look at a few additional examples which show typical biological processes and their SBGN Process Description representation. In Figure 2.2 a reversible reaction with two substrates and one product is shown. The enzyme E catalyzes an irreversible (metabolic) process which consumes two substrates (S1 and S2) 10

and produces one product (P1). The enzyme is a protein, therefore represented as a *macromolecule*. Substrates and product of the biochemical reaction are represented by *simple chemicals*. The consumption of S1 and S2 is represented by the *consumption arcs*. The *production arc* represents the synthesis of P1.

In Figure 2.3 the formation of a complex is shown. Two *macromolecule* entities X and Y form the *complex* X₂Y. Complex formation is represented using the *association* process node with ingoing *consumption* and outgoing *production* arcs. The *complex* glyph surrounds subunits X and Y.



Figure 2.2: This example of a Process Description shows an irreversible catalysis with 2 substrates and 1 products.




Figure 2.3: This example of a Process Description shows an irreversible catalysis with 2 substrates and 1 products.

In Figure 2.4 the regulation of a target gene by a transcription factor without knowledge about the promoter binding is shown. A transcription factor (TF) protein together with a target gene promoter X triggers the *process* of transcription. Direct binding of the TF to the target gene promoter has not been experimentally

verified, therefore the *logical operator* AND is used to describe the yet unspecified interaction between TF and target gene. The TF protein is a *macromolecule* of the *material type* 'protein' (mt:prot) whereas the gene promoter is given as a *nucleic acid feature* with the *conceptual type* 'gene' (ct:gene). The connecting arc *necessary stimulation* is applied to indicate that the stimulation by both regulator and target is necessary for the transcription process to take place. The target gene messenger as a product of the transcription process is represented by a *nucleic acid feature* with the *conceptual type* 'mRNA' (ct:mRNA). The *unspecified source* symbol is used to represent the large number of substrates of a transcription process (i.e. trinucleotides).

A last example is show in Figure 2.5, which shows passive transport or diffusion of a molecule. The *macromolecule* X in the cytosol serves as the substrate of a process leading to the production of the *macromolecule* X in the nucleus. This process describes the passive transport of X from one *compartment* to the other. The two macromolecules X do not carry the clone marker because the containing compartment is part of their identity.

More examples can be found in a list of so called SBGN bricks [?], which are building blocks representing basic biological patterns. These bricks can be used for assembly into different kinds of biological networks such as metabolic and regulatory networks.




Figure 2.4: This example of a Process Description shows a regulation of a target gene by a transcription factor without knowledge about the promoter binding.

2.2 Some more concept text

SBGN Process Description Level 1 is a visual language. Like any language like English it has a vocabulary, which in the case of SBGN is represented by the symbols that we call glyphs. Again like English our language has grammar, which are the underlying rules and concepts of the language that define its meaning. This specification aims to provide a detailed description of both the vocabulary and gram-



Figure 2.5: This example of a Process Description shows a passive transport of a molecule.

mar rules of the language as a reference. However, to understand this it is essential to grasp what an Process Description map describes.

The first thing to understand is the SBGN-PD does not describe biochemical reactions or gene expression — at least not directly. Instead this language describes collections — pools — of entities that are manipulated by processes which can convert them from one type of entity to another. An entity pool that is the input to a process is said to be consumed and that which is an output is produced. There is a third class of entity pool associated with a process: namely those that affect — modulate — the rate at which a process converts entity from one set of pools to another. The amount of entity in a modulating pools is not changed by the process. These very briefly are the key concepts of PD and are illustrated in figure ??.




Figure 2.6: A process consuming the entity pools A and B and producing the entity pools C and D. The entity pool E modulates the process. The process is represented as a square. Each (unspecified) entity pool is represented as an ellipsis identified by a label. The relationships between the process and each entity pool involved are represented by connecting lines called arcs. Depending on the type of relationship (consumption, production, modulation), these arcs are decorated with different symbols (respectively nothing, arrowhead, and diamond).

So how do this enable us to describe biological processes? Very simply. Figure ?? illustrates a biochemical reaction and shows how the essential components of that reaction, substrate, product and enzyme catalyst are all conveniently described by the representation. Likewise in figure ?? we can see again the concept of entity pools and process flow can be applied to gene regulation. This approach gives us a lot of flexibility and enables one to summarise biological processes that may incorporate many biochemical reactions (figure ??). 5

In the above examples of biological processes a number of different glyphs where used to illustrate different biological entities: the macromolecule, the simple chemical and the gene. 10

Graph, diagram or map?

A graph is a very technical term that belongs to mathematics and is uncommon in biology. Diagram is a concept that encompasses more than just graph. Examples are Venn diagrams for instance. Therefore, we recommend using the term map for SBGN representations. Those representations effectively permit users to travel and orient themselves in a biological system. Map is also the term most frequently used by the different communities, whether in metabolism, signaling or genomics. 15

Language Overview

3.1 Symbols used in SBGN Process Descriptions

An SBGN Process Description map is mainly a bipartite graph, i.e. it is made up of two types of nodes that connect in an alternate way (some exceptions are described below, e.g. when *logical operators* or *tag* are used). The two types of nodes are the *process nodes* and the *entity pools nodes*, the later representing the things that are modified by processes. These nodes are connected by arcs. In addition, the *entity pools nodes* can be contained in *compartments*.

Entity pool nodes

An entity pool is a population of entities that cannot be distinguished from each other, when it comes to the SBGN Process Description Level 1 map. For instance all the molecular entities that fulfill the same role in a given process form an entity pool. As a result, an entity pool can represent different granularity levels, such as all the proteins, all the instances of a given protein, only certain forms of a given protein. It really depends on what we want to represent. To belong to different compartments is sufficient to belong to different entity pools. Calcium ions in the endoplasmic reticulum and calcium ions in the cytosol belong to different entity pools when it comes to representing calcium release from the endoplasmic reticulum.

The Process Description language contains six glyphs representing classes of material entities: *unspecified entity* (Section 5.5), *simple chemical* (Section 5.5), *macromolecule* (Section 5.5), *nucleic acid feature* (Section 5.5), and *complex* (Section 5.5). (Specific types of macromolecules, such as protein, RNA, DNA, polysaccharide, and specific simple chemicals are not defined by Process Description but may be part of future levels of SBGN). In addition to the material entities, the Process Description language represents two conceptual entities: An absorbing pool, called *source and sink* (Section 5.5), and a *perturbing agent* (Section 5.5). Material

and conceptual entities can optionally carry auxiliary units such as *units of information* (Section 5.5), *state variables* (Section 5.5) and *clone markers* (Section 5.5).

Glyph: *Unspecified entity*

The simplest type of EPN is the *unspecified entity*: one whose type is unknown or simply not relevant to the purposes of the map. This arises, for example, when the existence of the entity has been inferred indirectly, or when the entity is merely a construct introduced for the needs of a map, without direct biological equivalent. For cases where the identity of the entities composing the pool is known, there exist other, more specific glyphs described below in the manual.

An *unspecified entity* is represented by an elliptic container, as shown in Figure 5.14. Note that this must remain an ellipse to avoid confusion with the Simple Chemical glyph, which is a circle (c.f. 5.5).




Figure 3.1: The Process Description glyph for *unspecified entity*.

Glyph: *Macromolecule*

Many biological processes involve *macromolecules*: biochemical substances that are built up from the covalent linking of pseudo-identical units. Examples of macromolecules include proteins, nucleic acids (RNA, DNA), and polysaccharides (glycogen, cellulose, starch, etc.). Attempting to define a separate glyph for all of these different molecules would lead to an explosion of symbols in SBGN, so instead, SBGN Process Description Level 1 defines only one glyph for all macromolecules. The same glyph is to be used for a protein, a nucleic acid, a complex sugar, and so on. The exact nature of a particular macromolecule in a map is then clarified using its label and decorations, as will become clear below. A *macromolecule* is represented by a rectangular container with rounded corners, as illustrated in Figure 5.17.




Figure 3.2: The Process Description glyph for *macromolecule*.

Examples of *macromolecules* are presented in Figure 3.3.




Figure 3.3: Examples of *macromolecules*. From left to right: the macromolecule of 18S ribosomal RNA, globin (a protein) in the oxygenated state, a molecule of starch (polymer of glucose), calcium calmodulin kinase 2 phosphorylated on threonine 286 and 306.

Glyph: *Simple chemical*

In SBGN Process Descriptions, a simple chemical is defined as the opposite of a macromolecule (Section 5.5): it is a chemical compound that is *not* formed by the covalent linking of pseudo-identical residues. Examples of simple chemicals are an atom, a monoatomic ion, a salt, a radical, a solid metal, a crystal, etc. A *simple chemical* is represented by a circular container, as depicted in Figure 5.12. To avoid confusion with the Unspecified Entity (5.5), this glyph must remain a circle and cannot be deformed into an ellipse. 5




Figure 3.4: The Process Description glyph for *simple chemical*.

Examples of *simple chemicals* are presented in Figure 3.5.




Figure 3.5: Examples of *simple chemicals*. From left to right: sodium chloride (a salt), dioxygen (a elemental molecule), calcium ion, glucose (an heteroatomics molecule), hydroxyl radical.

Glyph: *Nucleic acid feature*

10

The *Nucleic acid feature* construct in SBGN is meant to represent a fragment of a macromolecule carrying genetic information. A common use for this construct is to represent a gene or a transcript. The label of this EPN and its *units of information* (see Section 5.5) are often important for making the purpose clear to the

reader of a map. A *nucleic acid feature* is represented by a rectangular container whose bottom half has rounded corners, as shown in Figure 5.21. This design reminds that we are fundamentally dealing with a unit of information, but this information is carried by a macromolecule.




Figure 3.6: The Process Description glyph for *nucleic acid feature*.

Examples of *nucleic acid features* are presented in Figure 3.7.

5




Figure 3.7: Examples of *nucleic acid features*. From left to right: gene coding for the 18S ribosomal RNA, gene coding for α 1-globin, messenger RNA coding for α -globin, nuclear factor 1 binding site on the promoter of α 1-globin gene.

Glyph: **Complex**

A *complex* node represents a biochemical entity composed of other biochemical entities, whether macromolecules, simple chemicals, multimers, or other complexes. The resulting entity may have its own identity, properties and function in an SBGN map. A *complex* possesses its own container box surrounding the juxtaposed container boxes of its components. This container box is a rectangle with cut-corners (an octagonal box with sides of two different lengths). The size of the cut-corners are adjusted so that there is no overlap between the container and the components. The container boxes of the components must not overlap.

10




Figure 3.8: The Process Description glyph for *complex*.

Examples of *complexes* are presented in Figure 3.9.




Figure 3.9: Examples of *complexes*. From left to right: complex between phosphorylated cyclin and CDK2 forming the maturation promoting factor in yeast, a tetramer of complexes between globin and heme, and a complex between calcium-calmodulin kinase II and another complex, itself formed of calmodulin and calcium.

Glyph: *Empty Set*

It is useful to have the ability to represent the creation of an entity or a state from an unspecified source, that is, from something that one does not need or wish to make precise. For instance, in a model where the production of a protein is represented, it may not be desirable to represent all of the amino acids, sugars and other metabolites used, or the energy involved in the protein's creation. Similarly, we may not wish to bother representing the details of the destruction or decomposition of some biochemical species into a large number of more primitive entities, preferring instead to simply say that the species “disappears into a sink”. Yet another example is that one may need to represent an input (respectively, output)

into (resp. from) a compartment without explicitly representing a transport process from a source (resp. to a target).

For these and other situations, SBGN defines a glyph for explicitly representing the involvement of an unspecified source or sink. A *source* or *sink* is represented by the mathematical symbol for “empty set”, that is, a circle crossed by a bar linking the upper-right and lower-left corners of an invisible square drawn around the circle (\emptyset). Figure 5.10 illustrates this. Each source or sink node should be linked to one and only one arc in a map. The symbol used in SBGN is borrowed from the mathematical symbol for “empty set”, but it is important to note that it does not actually represent a true absence of everything or a physical void—it represents the absence of the corresponding structures in the model, that is, the fact that these sources or sinks are conceptually outside the scope of the map.




Figure 3.10: The *source* and *sink* glyphs.

Glyph: *Perturbing agent*

Biochemical networks can be affected by external influences. Those influences can be the effect of well-defined physical perturbing agents, such as a light pulse or a change in temperature; they can also be more complex and not well-defined phenomena, for instance the outcome of a biological process, an experimental setup, or a mutation. For these situations, SBGN provides the *perturbing agent* glyph. It is an EPN, and represents the amount to perturbing agent applied to a process. A *perturbing agent* is represented by a modified hexagon having two opposite concave faces, as illustrated in Figure 5.15.




Figure 3.11: The Process Description glyph for *perturbing agent*.

Glyph: *Multimer*

As its name implies, a multimer is an aggregation of multiple identical or pseudo-identical entities held together by non-covalent bonds. Thus, they are distinguished from polymers by the fact that the latter involve covalent bonds, and should be represented by *macromolecules*. Here *pseudo-identical* refers to the possibility that the entities differ chemically but retain some common global characteristic,

such as a structure or function, and so can be considered identical within the context of the SBGN Process Description. An example of this are the homologous subunits in a hetero-oligomeric receptor. SBGN Process Description accepts multimers of *simple chemical* (Section 5.5), *macromolecule* (Section 5.5), *nucleic acid feature* (Section 5.5) or *complex* (Section 5.5). A *multimer* is represented by two identical containers shifted horizontally and vertically and stacked one on top of the other as illustrated in Figure 3.12. 5




Figure 3.12: The Process Description glyph for *multimer*.

Examples of *multimers* are presented in Figure 3.13.




Figure 3.13: Examples of *multimers*. From left to right: pentameric nicotinic receptor in the open state, tetramer of oxygenated globin.

Decorations of the entity pool nodes

SBGN Process Description provides glyphs that decorate other glyphs, providing additional information that may be useful to the reader. These can provide annotation (*unit of information*), state information (*state variable*) or indicate duplication of entity pool nodes (*clone marker*). 10

Glyph: *Unit of information*

When representing biological entities, it is often necessary to convey some abstract information about the entity's function that is not related to its role in the map. The *unit of information* is a decoration that can be used in this situation 15

to add information to an EPN. Some example uses include: characterizing a logical part of an entity such as a functional domain (a binding domain, a catalytic site, a promoter, etc.), or the information encoded in the entity (an exon, an open reading frame, etc.). A *unit of information* can also convey information about the physical environment, or the specific type of biological entity it is decorating. A *unit of information* is represented by a rectangle overlapping the border of the EPN being annotated.

The label carried by *unit of information* defines the information it carries. For certain predefined types of information having controlled vocabularies associated with them, SBGN defines specific prefixes that must be included in the label to indicate the type of information in question. The controlled vocabularies predefined in SBGN Process Description Level 1 are described in Section 5.6.




Figure 3.14: The Process Description glyph for *unit of information*.

Glyph: *State variable*

Many biological entities, such as molecules, can exist in different *states*, meaning different physical or informational configurations. These states can arise for a variety of reasons. For example, macromolecules can be subject to post-synthesis modifications, wherein residues of the macromolecules (amino acids, nucleosides, or glucid residues) are modified through covalent linkage to other chemicals. Other examples of states are alternative conformations as in the closed/open/desensitized conformations of a transmembrane channel, and the active/inactive forms of an enzyme.

SBGN provides a means of associating one or more *state variables* with an entity; each such variable can be used to represent a dimension along which the state of the overall entity can vary. When an entity can exist in different states, the state of the whole entity (i.e., the SBGN object) can be described by the current values of all its *state variables*, and the values of the *state variables* of all its possible components, recursively. A *state variable* is represented by an elliptical container overlapping the border of the EPN being annotated.



Figure 3.15: The Process Description glyph for *state variable*.

A *state variable* does not necessarily have to be Boolean-valued. For example, an ion channel can possess several conductance states; a receptor can be inactive,

active and desensitized; and so on. As another example, a *state variable* “ubiquitin” could also carry numerical values corresponding to the number of ubiquitin molecules present in the tail. However, in all cases, a *state variable* on an EPN can only take *one* defined value. Further, an EPN’s *state variable* should always be displayed and always set to a value. An “empty” *state variable* is a *state variable* that is set to the value “unset”, it is not a *state variable* with no value. Note that the value “unset” is *not* synonymous to “any value” or “unknown value”.

Glyph: *Clone marker*

It is sometimes necessary to represent the same *EPN* several times. Otherwise, the resulting graph is so tightly connected that the map becomes unreadable. An example would be the representation of currency molecules such as ATP. However, we must indicate the fact, so that a reader knows the processes involving this particular glyph are not the only processes involving the *EPN*. If an *EPN* is duplicated on a map, we therefore mark all its graphical representation with a *clone marker* auxiliary unit. This marker provides the reader with a visual indication that this node has been cloned, and that at least one other occurrence of the *EPN* can be found in the map (or in a submap; see Section 5.5). The clone marker takes two forms, simple and labeled, depending on whether the node being cloned can carry state variables. Note that an *EPN* belongs to a single compartment. If two glyphs labelled “X” are located in two different compartments, such as ATP in cytosol and ATP in mitochondrial lumen, they represent different *EPNs*, and therefore do not need to be marked as cloned (and if they are, they are not part of the same clone).

The simple (unlabeled) *clone marker* is a portion of the surface of an *EPN* that has been modified visually through the use of a different shade, texture, or color. Figure 5.86 illustrates this. The *clone marker* occupies the lower part of the *EPN*. The filled area must be smaller than the unfilled one.



Figure 3.16: The Process Description glyph for *simple clone marker* applied to a *simple chemical*

Unlike the *simple clone marker*, the *labeled clone marker* includes (unsurprisingly, given its name) an identifying label that can be used to identify equivalent clones elsewhere in the map. This is particularly useful for stateful *EPNs*, because these can have a large number of state variables displayed and therefore may be difficult to visually identify as being identical. The filled area must be smaller than the unfilled one, but be large enough to have a height larger than the *clone marker*’s label (cf below).




Figure 3.17: The Process Description glyph for *labeled clone marker* applied to a *multimer of macromolecules*.

Glyph: *Clone marker*

It is sometimes necessary to represent the same *EPN* several times. Otherwise, the resulting graph is so tightly connected that the map becomes unreadable. An example would be the representation of currency molecules such as ATP. However, we must indicate the fact, so that a reader knows the processes involving this particular glyph are not the only processes involving the *EPN*. If an *EPN* is duplicated on a map, we therefore mark all its graphical representation with a *clone marker* auxiliary unit. This marker provides the reader with a visual indication that this node has been cloned, and that at least one other occurrence of the *EPN* can be found in the map (or in a submap; see Section 5.5). The clone marker takes two forms, simple and labeled, depending on whether the node being cloned can carry state variables. Note that an *EPN* belongs to a single compartment. If two glyphs labelled “X” are located in two different compartments, such as ATP in cytosol and ATP in mitochondrial lumen, they represent different *EPNs*, and therefore do not need to be marked as cloned (and if they are, they are not part of the same clone).

The simple (unlabeled) *clone marker* is a portion of the surface of an *EPN* that has been modified visually through the use of a different shade, texture, or color. Figure 5.86 illustrates this. The *clone marker* occupies the lower part of the *EPN*. The filled area must be smaller than the unfilled one.



Figure 3.18: The Process Description glyph for *simple clone marker* applied to a *simple chemical*

Unlike the *simple clone marker*, the *labeled clone marker* includes (unsurprisingly, given its name) an identifying label that can be used to identify equivalent clones elsewhere in the map. This is particularly useful for stateful *EPNs*, because these can have a large number of state variables displayed and therefore may be difficult to visually identify as being identical. The filled area must be smaller than the unfilled one, but be large enough to have a height larger than the *clone marker's* label (cf below).




Figure 3.19: The Process Description glyph for *labeled clone marker* applied to a *multimer of macromolecules*.

Controlled vocabularies used in SBGN Process Description Level 1

Some glyphs in SBGN Process Descriptions can be enriched with particular kinds of textual annotations conveying information relevant to the purpose of the glyph. These annotations are *units of information* (Section 5.5) or *state variables* (Section 5.5). An example is the case of multimers, which can carry a *unit of information* conveying the number of monomers composing the multimer. Other cases are described throughout the rest of this chapter.

In the rest of this section, we describe the controlled vocabularies (CVs) used in SBGN Process Description Level 1. They cover the following categories of information: an EPN's material type, an EPN's conceptual type, covalent modifications on macromolecules, the physical characteristics of compartments, and cardinality (e.g., of multimers). In each case, some CV terms are predefined by SBGN. With the exception of covalent modifications, the controlled vocabulary terms contained in *units of information* or *state variables* must be prefixed to indicate the type of information being expressed. Authors may use other CV values not listed here, but in such cases, they should explain the term's meanings in a figure legend or other text accompanying the map.

Entity pool node material types

The material type of an EPN indicates its chemical structure. A list of common material types is shown in Table 5.2, but others are possible. The values are to be taken from the Systems Biology Ontology ([?], <http://www.ebi.ac.uk/sbo/>), specifically from the branch *material entity* under *physical entity representation*. The labels are defined by SBGN Process Description Level 1.

Entity pool node conceptual types

An EPN's *conceptual type* indicates its function within the context of a given Process Description. In contrast to the *material types*, the *conceptual types* are not about physical composition, but about functional roles. For example, a strand of RNA is a physical artifact, but its use as messenger RNA is a role.

A list of common conceptual types is shown in Table 5.3, but others are possible. The values are to be taken from the Systems Biology Ontology (<http://>

3.1. Symbols used in SBGN Process Descriptions

Name	Label	SBO term
Non-macromolecular ion	mt:ion	SBO:0000327
Non-macromolecular radical	mt:rad	SBO:0000328
Ribonucleic acid	mt:rna	SBO:0000250
Deoxribonucleic acid	mt:dna	SBO:0000251
Protein	mt:prot	SBO:0000297
Polysaccharide	mt:psac	SBO:0000249

Table 3.1: A sample of values from the *material types* controlled vocabulary (Section 5.6).

www.ebi.ac.uk/sbo/), specifically from the branch *functional entity* under *physical entity representation*).

Name	Label	SBO term
Gene	ct:gene	SBO:0000243
Transcription start site	ct:tss	SBO:0000329
Gene coding region	ct:coding	SBO:0000335
Gene regulatory region	ct:grr	SBO:0000369
Messenger RNA	ct:mRNA	SBO:0000278

Table 3.2: A sample of values from the *conceptual types* vocabulary (Section 5.6).

Macromolecule covalent modifications

A common reason for the introduction of state variables (Section 5.5) on an entity is to allow access to the configuration of possible covalent modification sites on that entity. For instance, a macromolecule may have one or more sites where a phosphate group may be attached; this change in the site's configuration (i.e., being either phosphorylated or not) may factor into whether, and how, the entity can participate in different processes. Being able to describe such modifications in a consistent fashion is the motivation for the existence of SBGN's covalent modifications controlled vocabulary.

Table 5.4 lists a number of common types of covalent modifications. The most common values are defined by the Systems Biology Ontology in the branch *addition of a chemical group*, under *occurring entity representation*. The labels shown in Table 5.4 are defined by SBGN Process Description Level 1; for all other kinds of modifications not listed here, the author of a Process Description must create a new label (and should also describe the meaning of the label in a legend or text accompanying the map).

3.1. Symbols used in SBGN Process Descriptions

Name	Label	SBO term
Acetylation	Ac	SBO:0000215
Glycosylation	G	SBO:0000217
Hydroxylation	OH	SBO:0000233
Methylation	Me	SBO:0000214
Myristoylation	My	SBO:0000219
Palmytoylation	Pa	SBO:0000218
Phosphorylation	P	SBO:0000216
Prenylation	Pr	SBO:0000221
Protonation	H	SBO:0000212
Sulfation	S	SBO:0000220
Ubiquitination	Ub	SBO:0000224

Table 3.3: A sample of values from the *covalent modifications* vocabulary (Section 5.6).

Physical characteristics

SBGN Process Description Level 1 defines a special unit of information for describing certain common physical characteristics. Table 5.5 lists the particular values defined by SBGN Process Description Level 1.

Name	Label	SBO term
Temperature	pc:T	SBO:0000147
Voltage	pc:V	SBO:0000259
pH	pc:pH	SBO:0000304

Table 3.4: A sample of values from the *physical characteristics* vocabulary (Section 5.6).

Cardinality

5

SBGN Process Description Level 1 defines a special unit of information usable on multimers for describing the number of monomers composing the multimer. Table 3.5 shows the way in which the values must be written. Note that the value is an positive non-zero integer, and not (for example) a range. There is at present no provision in SBGN Process Description Level 1 for specifying a range in this context because it leads to problems of entity identifiability.

10

Process nodes

Process nodes represent processes that transform one or several entity pools into one or several entity pools, identical or different. SBGN Process Description Level 1

3.1. Symbols used in SBGN Process Descriptions

Name	Label	SBO term
cardinality	N : #	SBO : 0000364

Table 3.5: The format of the possible values for the *cardinality* unit of information (Section 3.1). Here, # stands for the number; for example, “N : 5”.

defines a generic *process* (Section 5.5), as well as five more specific ones: the *omitted process* (Section 5.5), the *uncertain process* (Section 5.5), the *association* (Section 5.5), the *dissociation* (Section 5.5), and the *phenotype* (Section 5.5).

Glyph: *Process*

A process is the basic process node in SBGN. It describes a process that transforms a given set of biochemical entities—macromolecules, simple chemicals or unspecified entities—into another set of biochemical entities. Such a transformation might imply modification of covalent bonds (conversion), modification of the relative position of constituents (conformational process) or movement from one compartment to another (translocation). A process transforms a set of entity pools (represented by *EPNs* in SBGN Process Description Level 1) into another set of entity pools. A *process* is represented by a square box linked to two connectors, small arcs attached to the centers of opposite sides. The consumption (Section 5.5) and production (Section 5.5) arcs are linked to the extremities of those connectors. The modulatory arcs (Section 3.1) point to the other two sides of the box. A *process* connected to *production* arcs on opposite sides is a reversible process.




Figure 3.20: The Process Description glyph for *process*.

The example in Figure 5.42 illustrates the use of a *process* node to represent a reaction between two reactants that generates three products. The stoichiometry for each entity pool involved is 1, and therefore can be omitted.




Figure 3.21: Reaction between ATP and fructose-6-phosphate to produce fructose-1,6-biphosphate, ADP and a proton.

The example in Figure 5.41 illustrates the use of a *process* node to represent the phosphorylation of a protein.




Figure 3.22: Phosphorylation of the protein MAP kinase.

The example in Figure 5.43 illustrates the use of a *process* node to represent a translocation. The large round-cornered rectangle represents a compartment border (see Section 5.5).

5




Figure 3.23: Translocation of calcium ion out of the endoplasmic reticulum. Note that the *process* does not have to be located on the boundary of the *compartment*. A *process* is not attached to any *compartment*.

The example in Figure 5.46 presents the conversion of two galactoses into a lactose. Galactoses are represented by only one *simple chemical*, the cardinality being carried by the *consumption* arc.




Figure 3.24: Conversion of two galactoses into a lactose.

Glyph: *Omitted process*

Omitted processes are processes that are known to exist, but are omitted from the map for the sake of clarity or parsimony. A single *omitted process* can represent any number of actual processes. For instance, one may want to represent a long chain of processes leading from one biochemical compound to another, without detailing all steps, but highlighting the fact that this is not a direct transformation. The *omitted process* is different from a *submap* (Section 5.5). While a *submap* references to an explicit content, that is hidden in the main map, the *omitted process* does not “hide” anything within the context of the map, and cannot be “unfolded”. An *omitted process* is represented by a *process* in which the square box contains a two parallel slanted lines oriented northwest-to-southeast and separated by an empty space. 5 10 15




Figure 3.25: The Process Description glyph for *omitted process*.

Glyph: *Uncertain process*

Uncertain processes are processes that may not exist. A single *uncertain process* can represent any number of actual processes. *Uncertain process* would be used to represent hypothesis, reactions which existence is supported by weak evidence etc. An *uncertain process* is represented by a *process* which square box contains a question mark. 20





Figure 3.26: The Process Description glyph for an *uncertain process*.

Glyph: Association

The association between one or more *EPNs* represents the non-covalent binding of the biological objects represented by those *EPNs* into a larger complex. An *association* between several entities is represented by a filled disc linked to two connectors separated by 180 degrees. The consumption (Section 5.5) and production (Section 5.5) arcs are linked to the extremities of those connectors. An *association* is never reversible, the inverse process being represented by a *dissociation* (Section 5.5).

Figure 3.27: The Process Description glyph for *association*.

The example in Figure 5.50 illustrates the association of cyclin and CDC2 kinase into the Maturation Promoting Factor.




Figure 3.28: Association of cyclin and CDC2 kinase into the Maturation Promoting Factor.

An *association* does not necessarily involve components of the same nature. Figure 5.51 gives an example illustrating the association of a pentameric *macromolecule* (a nicotinic acetylcholine receptor) with a *simple chemical* (the local anesthetic chlorpromazin) in an unnamed complex.




Figure 3.29: The association of a pentameric macromolecule with a simple chemical in an unnamed complex.

An association does not necessarily result in the formation of a *complex*; it can also produce a *multimer*. Figure 3.30 gives an example of using the successive formation of an hemoglobin monomer then a tetramer of the resulting complex.





Figure 3.30: Formation of hemoglobin.

Glyph: *Dissociation*

The dissociation of an *EPN* into one or more *EPNs* represents the rupture of a non-covalent binding between the biological entities represented by those *EPNs*. A *dissociation* between several entities is represented by two concentric circles. A simple empty disc could be, in some cases, confused with the *catalysis* (section Section 5.5). Moreover, the existence of two circles reminds the dissociation, by contrast with the filled disc of the *association* (Section 5.5).

Figure 3.31: The Process Description glyph for *dissociation*.**Glyph: *Phenotype***

A biochemical network can generate phenotypes or affect biological processes. Such processes can take place at different levels and are independent of the biochemical network itself. To represent these processes in a map, SBGN defines the *phenotype* glyph, which describes a process consuming nothing and producing nothing, but only modulated. A *phenotype* is represented by an hexagone, as illustrated in Figure 5.29. 5

Figure 3.32: The Process Description glyph for *phenotype*.

The example in Figure 3.33 illustrates the use of a *phenotype* node to represent cell division, stimulated by the mono-phosphorylated form of the maturation promoting factor (see Section 5.5 for the meaning of the open arrowhead). 10




Figure 3.33: Cell division stimulated by MPF.

Arcs

Arcs are lines that link nodes of SBGN together. The symbols attached to their extremities indicate their meaning. SBGN Process Description Level 1 defines nine arcs. *consumption* (Section 5.5), *production* (Section 5.5), *modulation* (Section 5.5), *stimulation* (Section 5.5), *catalysis* (Section 5.5), *inhibition* (Section 5.5), 15

and *necessary stimulation* (Section 5.5) connect *EPNs* to *PDs*. *LogicArc* (Section 5.5) link *EPNs* and *logic arcs*. *equivalenceArc* (Section 5.5) link nodes to *tag*. Arcs can take any shape, and are not restricted to segments of straight lines.

Glyph: *Consumption*

Consumption is the arc used to represent the fact that an entity pool is consumed by a process, but is not produced by the process. A *consumption* is represented by a simple line without particular symbols at its extremities. A cardinality label may be associated with *consumption* (Section 5.5) indicating the stoichiometry of the entity pool node for this process. This label is a number enclosed in a rectangle with one of the long sides adjacent to the consumption arc. Once assigned to one arc connecting to a process node, cardinality should be represented on all *consumption* and *production* arcs connected to that process node to avoid misinterpretation. In the case where the stoichiometry of some part of the process is not known, or undefined, a question mark (?) should be used within the cardinality label of the corresponding arcs.



Figure 3.34: The Process Description glyph for *consumption*.

Glyph: *Production*

Production is the arc used to represent the fact that an entity pool is produced by a process. In the case of a reversible process, the *production* arc also acts as a *consumption* arc. The target extremity of a *production* carries a filled arrowhead. A cardinality label can be associated with a *production* arc indicating the stoichiometry of a process.



Figure 3.35: The Process Description glyph for *production*.

Glyph: *Modulation*

A modulation affects the flux of a process represented by the target process. Such a modulation can affect the process **positively or negatively**, or even both ways depending on the conditions, for instance the concentration of the intervening participants. A *modulation* can also be used when one does not know the precise direction of the effect. The target extremity of a *modulation* carries an empty diamond.




Figure 3.36: The Process Description glyph for *modulation*.

Figure 5.74 represents the effect of nicotine on the process converting closed and open states of a nicotinic acetylcholine receptor. High concentrations of nicotine open the receptor while low concentrations can desensitize it without opening.




Figure 3.37: Modulation of nicotinic receptor opening by nicotine.

Glyph: *Stimulation*

5

A stimulation affects **positively** the flux of a process represented by the target process. This stimulation can be for instance a catalysis or a positive allosteric regulation. Note that *catalysis* exists independently in SBGN, see Section 5.5. The target extremity of a *stimulation* carries an empty arrowhead.




Figure 3.38: The Process Description glyph for *stimulation*.

The example in Figure 3.39 illustrates the use of two *stimulations* arcs to represent the opposite effects of agonists and inverse agonists on G-protein coupled receptor activity. Agonists stimulate the transition from inactive to active, while inverse agonists stimulate the transition inactive to active.

10




Figure 3.39: Opposite effects of agonists and inverse agonists on GPCRs.

Glyph: *Catalysis*

A catalysis is a particular case of stimulation, where the effector affects positively the flux of a process represented by the target process. The positive effect on the process is due to the lowering of the activation energy of a reaction. The target extremity of a *catalysis* carries an empty circle.

5




Figure 3.40: The Process Description glyph for *catalysis*.

The example in Figure 3.41 illustrates the use of *catalysis* arc to represent the effect of MAPKK on the phosphorylation of MAPK.




Figure 3.41: MAPKK catalyses the phosphorylation of MAPK.

Glyph: *Inhibition*

An inhibition **negatively** affects the flux of a process represented by the target process. This inhibition can be for instance a competitive inhibition or an allosteric inhibition. The target extremity of an *inhibition* carries a bar perpendicular to the arc.

5




Figure 3.42: The Process Description glyph for *inhibition*.

Glyph: *Necessary stimulation*

A necessary stimulation, is one that is necessary for a process to take place. A process modulated by a necessary stimulation can only occur when this necessary stimulation is active. The target extremity of a *necessary stimulation* carries an open arrow (to remind that it is a *stimulation*) coming after a larger vertical bar.

10




Figure 3.43: The Process Description glyph for *Necessary Stimulation*.

The example in Figure 5.79 below describes the transcription of a gene X, that is the creation of a messenger RNA X triggered by the gene X. The creation of the

protein X is then triggered by the mRNA X. (Note that the same example could be represented using the gene as reactant and product, although it is semantically different.)




Figure 3.44: The creation of a messenger RNA X triggered by the gene X.

The example in Figure 5.80 below describes the transport of calcium ions out of the endoplasmic reticulum. Without IP3 receptor, there is not calcium flux, therefore, one cannot use a *stimulation*. The Necessary Stimulation instead represents this absolute stimulation. 5




Figure 3.45: The transport of calcium ions out of the endoplasmic reticulum into the cytosol. Note that IP3R crosses both compartment boundaries. This is allowed, but the Macromolecule should only belong to one of the compartments.

Glyph: *Logic arc*

Logic arc is used to represent the fact that an entity influences the outcome of a logic operator. A *logic arc* is represented by a simple line without particular symbols at its extremities.




Figure 3.46: The Process Description glyph for *logic arc*.

Glyph: *Equivalence arc*

5

Equivalence arc is the arc used to represent the fact that all entities marked by a *tag* are equivalent. An *equivalence arc* is represented by a simple line without particular symbols at its extremities.




Figure 3.47: The Process Description glyph for *Equivalence arc*.

Logical operators

Glyph: *And*

10

The glyph *and* is used to denote that all the *EPNs* linked as input are necessary to produce the output. For instance a modulator A *and* a modulator B, when both present modulate the flux of a process. *And* is represented by a circle carrying the word “AND”.




Figure 3.48: The Process Description glyph for *and*. Only two inputs are represented, but more would be allowed.

Glyph: *Or*

15

The glyph *or* is used to denote that any of the *EPNs* linked as input is sufficient to produce the output. *Or* is represented by a circle carrying the word “OR”.




Figure 3.49: The Process Description glyph for *or*. Only two inputs are represented, but more would be allowed.

Glyph: *Not*

The glyph *not* is used to denote that the *EPN* linked as input cannot produce the output. *Not* is represented by a circle carrying the word “NOT”.




Figure 3.50: The Process Description glyph for *not*.

Glyph: *Compartment*

A compartment is a logical or physical structure that contains entity pool nodes. 5
 An *EPN* can only belong to one compartment. Therefore, the “same” biochemical
 species located in two different compartments are in fact two different “pools” and
 should be represented by two *EPNs*. A compartment is represented by a surface
 enclosed in a continuous border or located between continuous borders. These
 borders should be noticeably thicker than the borders of the *EPNs*. A compartment 10
 can take **any** geometry. A compartment must always be entirely enclosed.




Figure 3.51: The Process Description glyph for *compartment*.

To allow more aesthetically pleasing and understandable maps, compartments are allowed to overlap each other visually, but it must be kept in mind that this does not mean the top compartment contains part of the bottom compartment.

Glyph: Submap

A *submap* is used to encapsulate processes (including all types of nodes and edges) within one glyph. The *submap* hides its content to the users, and display only input terminals (or ports), linked to *EPNs* (Section 5.5). A *submap* is not equivalent to an *omitted process* (see Section 5.5). In the case of an SBGN description that is made available through a software tool, the content of a *submap* may be available to the tool. A user could then ask the tool to expand the *submap*, for instance by clicking on the icon representing the *submap*. The tool might then expand and show the *submap* within the same map (on the same canvas), or it might open it in a different canvas. In the case of an SBGN description made available in a book or a website, the content of the *submap* may be available on another page, possibly accessible via an hyperlink on the *submap*.

The *submap* is represented as a square box to remind the viewer that it is fundamentally a process. A *submap* carries labeled terminals. When the *submap* is represented folded, those terminals are linked to external *EPNs* (Section 5.5). In the unfolded view, exposing the internal structure of the *submap*, a set of *tags* point to the corresponding internal *EPNs* Section 5.5. A *tag* is represented by a rectangle fused to an empty arrowhead. The symbol should be linked to one and only one edge (i.e., it should reference only one EPN or compartment).




Figure 3.52: The Process Description glyph for *submap*. (Upper part) folded submap. (Lower part) content of the submap. The *uncertain process* represents the content that is not available outside the submap.

The left part of Figure 5.90 represents a *submap* that transforms glucose into fructose-6-phosphate. The *submap* carries five terminals, four linked to *EPNs* and one linked to a *compartment*. The latter is particularly important in the case of *EPNs* present only in a *compartment* enclosed in a *submap*, and that are not linked to terminals themselves. Note that the terminals do not define a “direction”, such as input or output. The flux of the reactions is determined by the context.

The map on the right of Figure 5.90 represents an unfolded version of the *submap*. Here, anything outside the *submap* has disappeared, and the internal *tags* are not linked to the corresponding external *terminals*. The yellow nodes are also present in the parent map, while the salmon nodes are specific to the submap.

Note the tag 5, linking the compartment “mito” of the *submap* to the compartment “mito” outside the *submap*. The compartment containing Glu6P is implicitly defined as the same as the compartment containing Glu and Fru6P. There is no ambiguity because if Glu and Fru6P were in different compartments, one of them should have been defined within the *submap*.

5

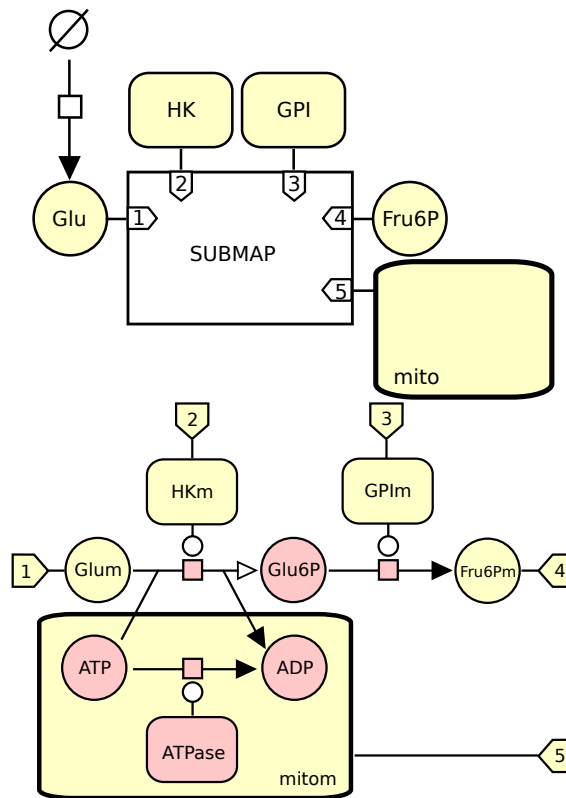


Figure 3.53: Example of a submap with contents elided.

3.2 Building a SBGN Process Description map

Now that the various symbols used by the SBGN Process Description language have been introduced, some guidance on building map will be provided.

How to choose the symbols to use?

It is important to realise that there are in general more than one way to represent a system in SBGN Process Description. The choice of concepts and symbols often depend on the granularity of information available, and the message the authors of the map wish to convey to the readers of the map.

As a first example of variable information granularity, let's take will take the example of MAP kinase phosphorylation (ERK). A very simple representation would

be to encode the state of phosphorylation in *entity pools node* with different names, ERK, ERK-P, ERK-PP (Figure 3.54).




Figure 3.54: Phosphorylation of ERK by MEK, where ERK is represented by three different macromolecules.

This kind of representation would be obtained if the SBGN Process Description map was generated from a model encoded in SBML core [?]. One of the problems with this representation is the difficulty for the reader to understand that the effect of MEK is to catalyse the phosphorylation of ERK. A scientist familiar with signalling pathways would probably immediately make the connection. A biologist in general could be more cautious. “P” could represent anything (peptide? proline?). A computer would require a special algorithm to parse the names, and this algorithm would easily fail. Instead of ERK-PP, we could have used PP-ERK, ERK_PP, ERKP1P2, ERKTPYP etc. But in fact, there is nothing in the the map Figure 3.54 indicating that the reactions catalysed add covalent modifications to ERK. Those reactions could be anything, such as aggregation, cleavage etc.

In order to overcome those issues, one can use state variables. On Figure 3.55, one uses a single state variables to represent the number of phosphorylations on ERK.




Figure 3.55: Phosphorylation of ERK by MEK, where ERK is represented by three different states, non-phosphorylated, mono-phosphorylated and bi-phosphorylated.

Because ‘P’ is a reserved symbol of the covalent modifications vocabulary, there is no ambiguities. We know that each reaction add a phosphate to ERK. That would

be the representation to favour if we only know the number of phosphorylation (e.g. by western blot with non-specific antibodies), or if we do not care which site is phosphorylated. Note that the leftmost ERK carries an empty state variable, that is equivalent to “0P”. The state variable is not omitted. This rule of SBGN Process Description Level 1 is called “once a variable, always a variable” (OVAV). If we want to, or can, be more specific about distinct phosphorylations, one can create two state variables, one for each site (Figure 3.56). 5




Figure 3.56: Phosphorylation of ERK by MEK, where each phosphorylated form of ERK is represented.

In this representation, we have all the information related to the two phosphorylation sites, the threonine and the tyrosine. They are represented by the variable symbols T and Y in the figure. But one could have chosen X and Y or 1 and 2. The important issue is to distinguish them. Note that the creation of an extra *entity pool node* is unavoidable. SBGN Process Description Level 1 does not currently allow logical expressions in the state variables. Therefore if only one *entity pool node* was to be used to represent the single-phosphorylated form, a choice between T or Y should have been made, and the resulting map would not have carried the same information than (Figure 3.55). 10 15

As a second example, we will consider the oxygenation of hemoglobin. If one wants to convey the message that 4 oxygen molecules bind to a molecule of hemoglobin, it is sufficient to create *macromolecules* for hemoglobin and oxy-hemoglobin.




Figure 3.57: Hemoglobin oxygenation using *macromolecules*.

If conveying the fact that hemoglobin is a multimer is important (for instance to suggest cooperativity), one can use *multimers* instead. In addition, in Figure 3.58, the concept of oxygenation is represented with a state variable rather than being embedded in the name of the nodes.




Figure 3.58: legend

An additional layer of complexity is needed if we want to mention the α and β subunits. A *complex* can then be used. In addition Figure 3.59 explicitly represent the complexes between globin and oxygene instead of using state variables.




Figure 3.59: legend

In conclusion, one can see that many different choice are offered to represent an idea in SBGN, and the map writers make the choice. By doing so, they acknowledge that more or less information can be extracted from the resulting map. The important issue is that anyone reading the map interpret it the same way. That is the topic of the following section.

5

Worked Examples

Part II

Technical Reference

Language Elements

5.1 Introduction

In this chapter we aim to describe the Process Description language by describing its elements and underlying concepts in detail. The challenge in doing this is to provide enough detail to minimise ambiguity, but to also make the rules understandable to the users of the specification. An additional goal is to minimise duplication of rules as much as possible, which makes the specification more difficult to update in the future. To achieve this we took our lead from other successful standards: the Unified Modelling Language (UML) [?] which is a graphical language that has been used by tens of thousands of software developers for over a decade to describe a software system; and the Systems Biology Markup Language (SBML) Level 3 Core specification [?] which is the “original” computational systems biological standard and after 10 years is firmly embedded in the scientific community it serves. In both these specifications the language is modelled in UML where each language element is described as an UML class. The specification defines each class in turn and in so doing describes how the language elements fit together (syntax) and how they are applied and their rules (semantics). Since UML is a graphical language, the class definition also includes, where applicable to the language element, a description of the symbols used and guidelines about how the symbols should be drawn or laid out.

Based on these specifications we will follow the following conventions in this chapter:

- The language will be described by a UML class model, a language element of language concept being represented as a class.
- Each class will be defined in detail. The definition will describe parent classes, attributes, interactions with other classes and any rules that apply to that class.

- Where classes correspond to glyphs, or where appropriate, a detailed description of the glyph will be given.
- Any rules or concepts that apply to the language as a whole are defined after the individual class definitions.

5.2 Note on typographical convention

5

The concept represented by a glyph is written using a normal font, while a *glyph* means the SBGN visual representation of the concept. For instance “a biological process is encoded by the SBGN PD *process*”. A UML class name is written in camel case and presented as `CamelCaseClassName`. Attributes and associations are written in lower case, using the underscore to separate words: `an_attribute`.

10



See footnote.

Note that in drafts of this document the margin-note¹ is used to indicate that this part of the document may require special attention from the reader as it may contain material that is new or potentially contentious.

5.3 How to read the Language Specification

Here will describe the elements of the class definitions that are the core of the language specification. Each definition starts with an introduction describing the purpose of the class and what it represents conceptually or physically. It's context within the UML Process Description language is then described by a detailed figure showing the class, its associations and its interaction with other key classes to help understand its context and any rules that are part of the definition.

15

20

A number of terms are used within the specification that for clarity we define here:

class A class describes a set of objects that share the same specifications of features, constraints, and semantics (from the UML specification [?]).

subclass A class that inherits attributes, behaviour and associations from another class. For example in figure 5.1 “class C is a subclass of class A”.

25

superclass A class that is an ancestor of another class. For example in figure 5.1 “class A is a superclass of class B”.

generalisation A relationship in UML that defines a subclass/superclass relationship.

30

association A semantic relationship between two classes. Typically at least one, but often both classes require the other class to complete its definition.

¹The footnote contains a description of why this part of the specification is contentious, may require further discussion, or should be noted by the reviewer of the document.

instance An instance is equivalent to an object — a single example or realisation of a class.

role The role describes the nature of an association from the perspective of one of the classes in that relationship.

cardinality In the case of an attribute the number of separate values the attribute can hold and for an association the number of instances of each class that can be associated with each other. The permitted cardinality values are:

R Required (attributes only).

O Optional (attributes only).

1 required - only one instance is permitted (association only). 10

0..1 optional - zero or one instance is permitted (association only).

1..* at least 1 instance is required.

***** any number of instances are permitted.

type A type constrains the values represented by an attribute. For example an attribute of type `int` must have a value that is a integer. 15

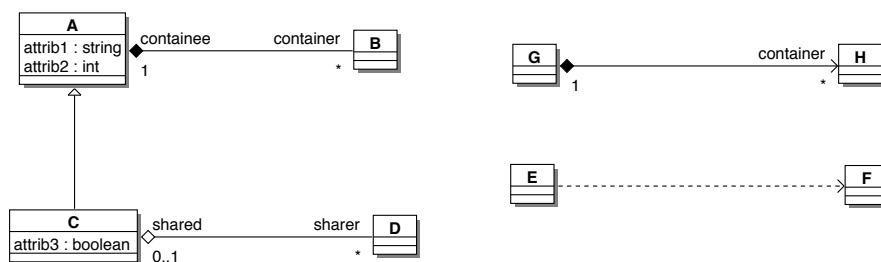


Figure 5.1: An example of the UML elements used in this specification. The diagram is explained more fully in the text.

The UML class diagrams used will look something like the examples in figure 5.1. The UML elements in this diagram are the only ones used in this specification and again as an aid to understanding their meanings are described below:

class See above for definition. The class symbol displays each attribute and its type using the convention: “attribute : type”. 20

generalisation (open triangle arrowhead) Defines an inheritance relationship where class C inherits the attributes, associations and semantics of class A. This means that C also has attributes `attrib1` and `attrib2` as well as an association called `container` to class B.

composite aggregation (black diamond) The aggregating class (A) is the one adjacent to the diamond and at the other end is the contained class (B). This represents a whole/part relationship where class B is part of class A and B cannot exist independently of A. The cardinality and role of a given class are shown at the opposite end of the association to it. 5

shared aggregation (open diamond) The aggregating class (C) is adjacent to the diamond and is connected to the other class (D) via a solid line. Here there is a whole/part relationship, but there the part (class D) can be shared with another class and can exist independently of C.

navigability (black diamond anchor arrowhead) If the association ends in an arrow, this indicates the direction of the association. Here this should be taken to mean that only one side of the association (G) is 'aware' of the relationship. 10

dependency Indicates that the dependent class (E) requires the other class adjacent to the arrow head (F) to satisfy its specification or implementation. 15

The specification uses a number of primitive types that are used in attribute definitions. These are:

int An integer.

string A string of UniCode characters.

boolean A Boolean value that can be either True or False. 20

object A type that can be any value.

cv A controlled vocabulary (see 5.6).

enum A value that must be chosen from one of an enumerated set of predefined values.

An example definition 25

Here we will provide an example specification. In this section is an overview of the class and its concepts.

Generalisation

This section defines the inheritance relationship(s) between this class and any other classes. 30

Attributes

Here any attributes specific to this class are defined and their meaning or purpose described. Attributes from superclasses are part of this classes definition but are not defined here explicitly.

Associations

5

Any associations are defined and described. Again associations from superclasses are not included here, but are part of this classes definition.

Notation

If the class corresponds to one or more glyphs then each glyph is described here. The glyph is described in words and graphically. In cases where several glyphs combine in complex ways usage examples are provided. 10

Layout Rules and Guidelines

In some cases the graphical layout of a glyph and is more complicated and requires some additional explanation. If this is the case this will be provided here.

Rules and Constraints

15

The semantics of the class are defined here in the form of itemised rules and constraints on the behaviour of the class. The scope of this section is ideally restricted to rules and behaviours the related to the class itself or classes it has some form of immediate relationship with (an association or dependency). In the later case the rules should relate to that relationship. 20

Changes from Previous Version

In order to help track changes between versions of the specification this section documents where this class definition differs from that in previous versions. Where appropriate ticket numbers for bugs or issues addressed in this version should be included. 25

5.4 Overview

The UML model describing the Process Description language is summarised in figure 5.2. The model has a root class `SBGNElement` for all language elements and care has been taken keep the model as simple as possible. This includes minimising the use of multiple inheritance and reflecting the directed graph structure inherent in the language with the `SBGNNode` and `SBGNArc` classes. All graphical elements that can be drawn directly onto a Process Description map are glyphs (`SBGNGlyph`) and all those that decorate glyphs are auxiliary units 30

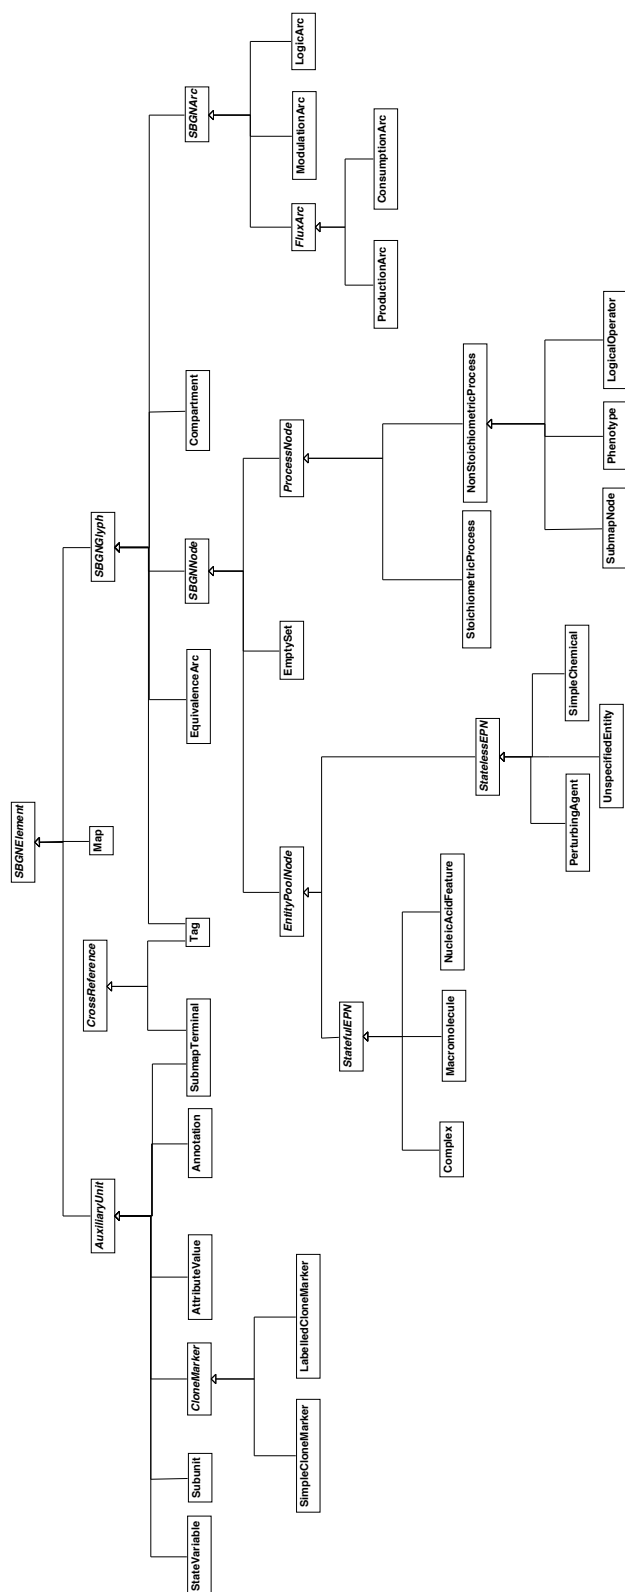


Figure 5.2: A view of the UML model describing Process Description language. This diagram shows the classes and their inheritance relationships. No attributes or associations between classes are shown to simplify the diagram. These details are provided in the UML class diagrams associated with most class definitions.

(AuxiliaryUnit). There are a number of language rules that the model and the individual class specification cannot capture and these are dealt with later in the chapter after the class definitions in section 5.5.

5.5 Definitions

SBGNElement 5

⚠ See footnote.

All the glyphs in SBGN Process Description Level 1 inherit from SBGNElement². This is an abstract or conceptual class that helps organise Process Description conceptually. SBGNElement (figure 5.3) has a single attribute id that is an identifying attribute. This means that all SBGN elements defined here, which ultimately extend SBGNBase, can all be uniquely identified from each other. This makes sense if you think that a glyph drawn on a map is distinct from another glyph drawn on the map. The id attribute reflects this and is not shown explicitly in a Process Description map. 10

Generalisation None 15

Attributes

⚠ See footnote.

id: identifier (R) uniquely identifies all SBGN elements in the same namespace³.

Changes from Previous Version Not defined in the previous version.

Map 20




Figure 5.3: UML definition of Map and SBGNElement.



See footnote.

The Map⁴ (figure 5.3) is a container that holds all the glyphs (SBGNGlyph (see section 5.5)) drawn in a Process Description map. A map may represent a submap or a supermap and should comply with the rules set out in section 5.8. The elements held should be logically unique and conform to the identity rules in section 5.7.

5

Generalisation

- SBGNElement (see section 5.5)

Attributes

No additional attributes

Associations

10

elements:SBGNGlyph (*) The collection of glyphs held by the map.

Rules and Constraints

- A map is valid if it is empty (although not very useful).
- All instances of SBGNGlyph (see section 5.5) must be unique (see section 5.7).

15

Notation

The map is the canvas upon which the Process Description language is drawn. It's only visible feature is its colour. It can take any pattern or colour (or be transparent for that matter), but as SBGN is 'colour blind' this does not convey any meaning in itself.

20

Changes from Previous Version

Not defined explicitly in previous versions.

SBGNGlyph

The SBGNGlyph is the fundamental building blocks of the Process Description language. It is the only element that can be drawn directly on a map (Map (see section 5.5)).

25

²A new concept

³Not defined previously, but doesn't change Process Description semantics. reinforces the idea of instance identity that exists for all glyphs.

⁴Defining this explicitly is new, but the concept of the map as a container of glyphs has always existed in the Process Description language




Figure 5.4: UML definition of the SBGNGlyph and its subclasses.

Generalisation

- SBGNElement (see section 5.5)

Attributes

No additional attributes.

Associations

5

owning_map:Map (1) The map that contains this class.

Rules and Constraints

No additional rules and constraints.

Changes from Previous Version

Not defined in previous version.

10

AuxiliaryUnit

The AuxiliaryUnit (figure 5.5) represents symbols that may be used to adorn glyphs. In doing so they change the meaning of the glyph and/or provide additional information about it.

Generalisation

15

- SBGNElement (see section 5.5)

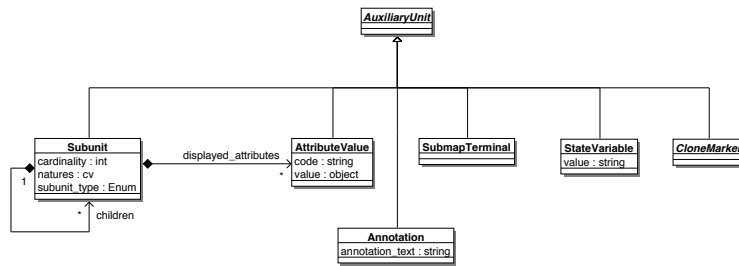


Figure 5.5: UML definition of the Auxiliary Unit and its subclasses.

Attributes

No additional attributes.

Associations

No additional associations.

Rules and Constraints

5

No additional rules and constraints.

Changes from Previous Version

Not defined in previous version.


SBGNNode

Figure 5.6: UML definition of the SBGNNode and SBGNArc classes.

The **SBGNNode** (figure 5.6) represents the nodes in the graph structure that is the core representation within the Process Description language. The nodes are connected to glyphs descended from **SBGNArc** for form a direct graph. 10

Generalisation

- **SBGNGlyph** (see section 5.5)

Attributes

No additional attributes.

Associations

out_arcs:SBGNArc (*) The arcs that are leaving this node.

in_arcs:SBGNArc (*) The arcs that are entering this node.

5

Rules and Constraints

- The set of SBGNNodes linked to this node via a SBGNArc (its adjacent nodes) must be all belong to different entity pools (as defined by Entity-PoolNode) and cannot include more than one clone of the same entity pool.

Changes from Previous Version

10

Not defined in the previous version.

SBGNArc

The SBGNArc (figure 5.6) represents the directed arcs (also know as directed edges) in the directed graph structure that is the core representation within Process Description language. The arc is connected to two nodes descended from SBGNNode, one at each end. As the arc has a direction these nodes are by convention designated the out_node to indicate the nodes that the arc is leaving and in_node to indicate the node that it is entering.

15

Generalisation

- SBGNGlyph (see section 5.5)

20

Attributes

No additional attributes.

Associations

out_node:SBGNNode (1) The node that this arc is leaving.

in_node:SBGNNode (1) The node that this arc is entering.

25

Rules and Constraints

No additional rules and constraints.

Changes from Previous Version

Not defined in the previous version.

EntityType




Figure 5.7: UML definition of the entity type and the state variable definition. The diagram shows how these classes interact with the entity pool, state variable and so influence EntityPoolNode logical identity.



See footnote.

The **EntityType**⁵ is one of the core concepts in the Process Description language. It defines the type of entity that is instantiated by one or more entity pools in a Process Description map. The **EntityType** has associated state variable definitions (see figure 5.7) and this enforces one of the core rules in the Process Description language that once a state is associated with an entity type it must be used by all entity pools of that type. 5

Generalisation

None.

Attributes

10

name: string (R) The name that identifies the entity in the Process Description map. EPNs with the same label should be from the same entity. The string cannot be empty and must start and end with a non-space character. Any Unicode character is acceptable⁶.



See footnote.

natures: cv(*) The nature of the entity pool node as defined by a controlled vocabulary. Zero, one or more values may be set, but each one must belong to a different controlled vocabulary (see section 5.6)⁷. 15



See footnote.

Associations

state_definitions:StateVariableDefintion (*) The state definitions associated with this type. 20

⁵The concept of an entity pool's type has been there implicitly in previous versions and comes up in discussions. This class aims to formalise that concept and the rules associated with it and enable us to formalise rules associated with EPNs.

⁶Not discussed or defined anywhere, but would make sense to define this explicitly.

⁷This has been discussed on the mailing list where this seems to be the consensus solution

pools:EntityPoolNode (1..*) The entity pool nodes that used this type.

Logical Identity

Logical Key:

- name
- natures

5

Rule and Constraints

- All instances of **EntityPoolNode** associated with a particular **EntityType** must be of the same class.
- If an instance of **EntityType** contains one or more instances of **StateVariableDefinition** then the **EntityPoolNodes** associated with it must be subclasses of **StatefulEPN**.

10

Changes from Previous Version

Although not defined explicitly in the previous version, this concept and the associated rules did exist in the language.

StateVariableDefinition

15



See footnote.

The **StateVariableDefinition**⁸ defines the state variables used by an **EntityType** and therefore those state variables that must exist in an **EntityPoolNode** (see figure 5.7).

Generalisation

None.

20

Attributes

name: string (O) The name that of the state variable. This is optional, but if defined cannot be an empty string or just white space characters. It should also start with an alpha-numeric character and end with a non-space character. It should not contain a '@' character⁹.

25



See footnote.

Associations

owning_entity_type:EntityType (1) The **EntityType** that owns this definition.

⁸As with **EntityType** this is new and aims to formalise the concept that an entity pool must preserve the same state variables whenever it is used.

⁹No rule defined previously, but this would seem to make sense.

Rule and Constraints

None.

Changes from Previous Version

Although not defined explicitly in the previous version, arguably this concept did exist in the language.

5

EntityPoolNode

An entity pool is a population of entities that cannot be distinguished from each other, when it comes to the SBGN Process Description Level 1 map. For instance all the molecular entities that fulfill the same role in a given process form an entity pool. As a result, an entity pool can represent different granularity levels, such as all the proteins, all the instances of a given protein, only certain forms of a given protein. To belong to a different compartment is sufficient to belong to different entity pools. Calcium ions in the endoplasmic reticulum and calcium ions in the cytosol belong to different entity pools when it comes to representing calcium release from the endoplasmic reticulum.

10

15

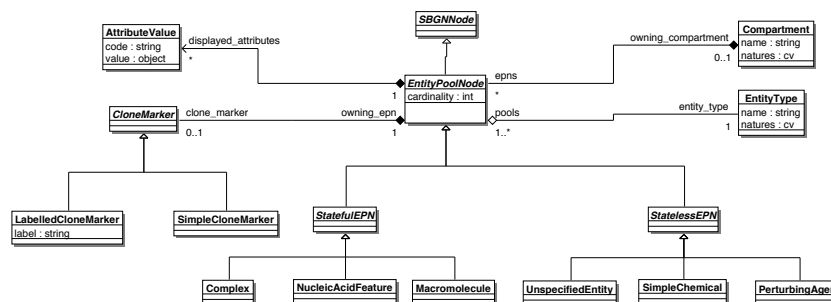


Figure 5.8: UML definition of the entity pool node and its descendant glyphs.



See footnote.

The EntityPoolNode (figure 5.8) is the definition of the entity pool and it shares an EntityType with other identical entities¹⁰. An instance of an entity pools is therefore distinguished from other pools with the same entity type by its cardinality, its owning_compartment and the values of its StateVariables (where appropriate). It must belong to a compartment or be associated with the map (c.f. section 5.5) and can contain a clone marker if it is cloned (see section 5.7)—note that not all EPNs can be cloned.

20

Generalisation

- SBGNNode (see section 5.5)

¹⁰Although this concept is discussed it is not explicitly defined previously.

Attributes

cardinality: int (R) The number of copies of the entity. Must be a positive non-zero integer.

Associations

owning_compartment:Compartment (0..1) The compartment that this EPN belongs too. 5

entity_type:EntityType (1) The type of this entity pool.

clone_marker:CloneMarker (0..1) The clone marker decorator. See section 5.5 for its use.

displayed_attributes:AttributeValue (*) One or more decorators used to display attribute values¹¹. 10



See footnote.

Logical Identity

Logical Key:

- **owning_compartment**
- **entity_type** 15
- **cardinality**

Rules and Constraints

- If **cardinality** > 1 then the descendant glyph must be displayed as a multi-mer.
- If the EPN is drawn directly on a *Map* then **owning_compartment** is not set. We interpret this as belonging to an invisible default compartment. 20
- **natures** can only use the material type (section 5.6), conceptual type (section 5.6) or physical characteristics (section 5.6) controlled vocabularies.
- The appropriate subclass of **CloneMarker** must be used to distinguish logically identical instances of this class. 25
- the **EntityPoolNode** must be associated with at least one **SBGNArc** (see section 5.5) (degree > 0).
- All **StateVariableDefinitions** associated with the **EntityType** must have an associated **StateVariable**.

¹¹This is an alternate way of using the Unit of Information to display information, but to constrain it so that it presents attributes of the EPN not general annotation. See the **AttributeValue** class for more information.

Notation

Although there is no direct graphical representation of this class the appearance of the `AttributeValue` and its associated glyph the *Unit of Information* is common to all subclasses so it is convenient to describe it here. The `AttributeValue` can be used to present the cardinality and natures of an EPN subclasses. These use the following codes to indicate which attribute is being presented:

<code>pc</code>	container physical characteristic	
<code>mt</code>	entity pool material type	
<code>ct</code>	entity pool conceptual type	
<code>N</code>	multimer cardinality	10

Changes from Previous Version

Not defined explicitly in the previous version, but the concept of the EPN and its semantics existed. The main change to previous semantics is that of the natures, which didn't formally exist before, but which now must contain a unique set of controlled vocabularies and is part of the logical key of the `EntityPoolNode`.

Empty Set

It is useful to have the ability to represent the creation of an entity or a state from an unspecified source, that is, from something that one does not need or wish to make precise. For instance, in a model where the production of a protein is represented, it may not be desirable to represent all of the amino acids, sugars and other metabolites used, or the energy involved in the protein's creation. Similarly, we may not wish to bother representing the details of the destruction or decomposition of some biochemical species into a large number of more primitive entities, preferring instead to simply say that the species "disappears into a sink". Yet another example is that one may need to represent an input (respectively, output) into (resp. from) a compartment without explicitly representing a transport process from a source (resp. to a target).

For these and other situations, SBGN defines a single glyph to handle these situations representing the involvement of an external pool of entities. The symbol used in SBGN is borrowed from the mathematical symbol for "empty set", but it is important to note that it does not actually represent a true absence of everything or a physical void—it represents the absence of the corresponding structures in the model, that is, the fact that the external pool is conceptually outside the scope of the map.

A frequently asked question is, why bother having an explicit symbol at all? The reason is that one cannot simply use an arc that does not terminate on a node, because the dangling end could be mistaken to be pointing to another node in the map. This is specially true if the map is rescaled, causing the spacing of elements

in the map to change. The availability and use of an explicit symbol for sources and sinks is critical.

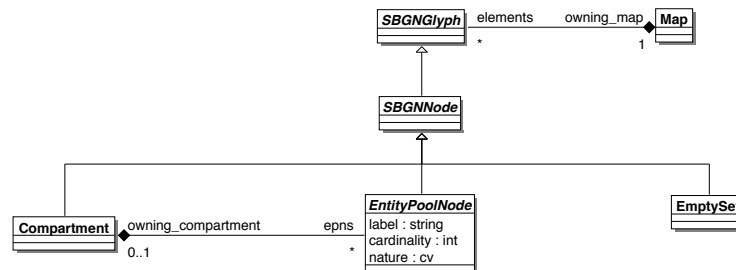


Figure 5.9: The UML definition of the *EmptySet* and its context in relation to other elements of the Process Description language.

The definition of the *Empty Set* is shown in figure 5.9. The empty set is not a subclass of *EntityPoolNode* as it does not represent a single pool of entities and does not share any of the other attributes of an *EntityPoolNode*, nor does it belong to a particular compartment¹². 5

⚠
See footnote.

Generalisation

- *SBGNNode* (see section 5.5)

Attributes

No additional attributes. 10

Associations

No additional associations.

Rules and Constraints

- All instances of *Empty Set* can be regarded as identical therefore not special decoration is used to indicate replication on the map. 15
- the *EmptySet* must be associated with at least one *SBGNArc* (see section 5.5) (degree > 0).

Notation

Glyph: *Empty Set*

SBO Term: SBO:0000291 ! empty set 20

¹²This is a significant change to the semantics from v1.3 since it is no longer an EPN.

Container: Represented by the mathematical symbol for “empty set”, that is, a circle crossed by a bar linking the upper-right and lower-left corners of an invisible square drawn around the circle (\emptyset). Figure 5.10 illustrates this. The symbol should be linked to one and only one edge in a map.

Label: None

5




Figure 5.10: The *empty set* glyph.

Changes from Previous Version

The EmptySet and *Empty Set* glyph has replaced the *Source* and *Sink* glyphs. This symbols used remains the same, but the underlying concept has changed. The *Source* and *Sink* glyphs where types of EPN, representing single entity pools, while the EmptySet is not.

10

StatelessEPN

The StatelessEPN (figure 5.11) represents a pool where the entities do not change ‘state’. In other-words the entities do not undergo any physical change that is useful to record in a Process Description map. Therefore they cannot be assigned a state-variable.

15

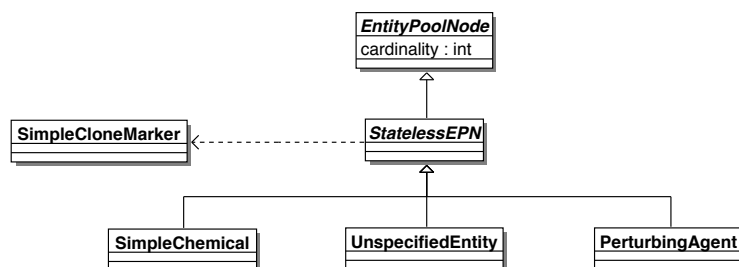


Figure 5.11: UML definition of the stateless entity pool node and its descendant glyphs.

Generalisation

- EntityPoolNode (see section 5.5)

Attributes

No additional attributes.

Associations

No additional associations.

Rules and Constraints

- if a clone marker is used it must be of type SimpleCloneMarker.

Changes from Previous Version

5

Not defined in the previous version.

Simple chemical

A SimpleChemical is the ‘opposite’ of a macromolecule (Section 5.5): it is a chemical compound that is *not* formed by the covalent linking of pseudo-identical residues. Examples of simple chemicals are an atom, a monoatomic ion, a salt, a radical, a solid metal, a crystal, etc.

10

Generalisation

- StatelessEPN (see section 5.5)

Attributes

No additional attributes.

15

Associations

No additional associations.

Rules and Constraints

No additional rules and constraints.

Notation

20

There are two glyphs associated with SimpleChemical. The first *simple chemical monomer* is used when cardinality = 1 and the second *simple chemical multimer* is used when cardinality > 1.

Glyph: Simple chemical monomer

SBO Term: SBO:0000247 ! simple chemical

25

Container: A *simple chemical* is represented by a ‘stadium’ symbol: a circle split in two with a rectangle inserted between them (see figure 5.12). If desired the rectangle can have zero length and the symbol is then identical to a circle

(Figure 5.12). To avoid confusion with the Unspecified Entity (5.5), this form of the glyph must remain a circle and cannot be deformed into an eclipse.

Label: The identification of the *simple chemical* is carried by an unbordered box containing a string of characters. The characters may be distributed on several lines to improve readability, although this is not mandatory. The label box has to be attached to the center of the circular container. The label is permitted to spill outside the container.



Figure 5.12: The Process Description glyph for *simple chemical monomer*. The stadium form and circular forms are shown, as are the cloned forms of the glyph.

Glyph: *Simple chemical multimer*

SBO Term: SBO:0000421 ! multimer of simple chemicals

Container: A *simple chemical multimer* is represented by two identical containers shifted horizontally and vertically and stacked one on top of the other. Figure 5.13 illustrates the glyph.

Label: The multimer carries an identifying label. The label is placed in an unbordered box containing a string of characters. The characters can be distributed on several lines to improve readability, although this is not mandatory. The label box must be attached to the center of the top monomer's container. The label may spill outside of the container.




Figure 5.13: The Process Description glyph for *simple chemical multimer*. The figures show the stadium and circular forms, and their cloned variants.

Changes from Previous Version

The glyphs used for the SimpleChemical have been changed to the stadium glyph. Previously the glyph was a circle. To maintain compatibility with previous versions the stadium symbol can be drawn without straight horizontal elements so that it becomes a circle.

UnspecifiedEntity

The simplest type of EntityPoolNode is the UnspecifiedEntity — one whose type is unknown or simply not relevant to the purposes of the map. This arises, for example, when the existence of the entity has been inferred indirectly, or when the entity is merely a construct introduced for the needs of a map, without direct biological relevance. These are examples of situations where the UnspecifiedEntity is appropriate. (Conversely, for cases where the identity of the entities composing the pool is known, there exist other, more specific glyphs described elsewhere in the specification.)

Generalisation

- StatelessEPN (see section 5.5)

Attributes

No additional attributes.

Associations

No additional associations.

Rules and Constraints

- The UnspecifiedEntity cannot have cardinality > 1. This means there is no multimer glyph.
- It cannot have a *natures*¹³, therefore *natures* must **not** be set.



See footnote.

Notation

Glyph: *Unspecified entity*

SBO Term: SBO:0000285 ! material entity of unspecified nature

Container: An *unspecified entity* is represented by an elliptic container, as shown in 5.14. Note that this must remain an ellipse to avoid confusion with the Simple Chemical glyph, which is a circle (c.f. 5.5).

Label: An *unspecified entity* is identified by a label placed in an unbordered box containing a string of characters. The characters can be distributed on several lines to improve readability, although this is not mandatory. The label box must be attached to the center of the container. The label may spill outside of the container.

¹³Not explicitly discussed, but *natures* don't make sense here.



Figure 5.14: The Process Description glyph for *unspecified entity*.

Changes from Previous Version

No changes from the previous version.

Perturbing Agent

Biochemical networks can be affected by external influences. Those influences can be the effect of well-defined physical perturbing agents, such as a light pulse or a change in temperature; they can also be more complex and not well-defined phenomena, for instance the outcome of a biological process, an experimental setup, or a mutation. For these situations, SBGN provides the *perturbing agent* glyph. It is an EPN, and represents the amount to perturbing agent applied to a process. 5 10

Generalisation

- StatelessEPN (see section 5.5)

Attributes

No additional attributes.

Associations

No additional attributes. 15

Rules and Constraints

- The PerturbingAgent cannot have cardinality > 1 . This means there is no multimer glyph.

Notation 20

Glyph: *Perturbing agent*

SBO Term: SBO:0000405 ! perturbing agent

Container: A *perturbing agent* is represented by a modified hexagon having two opposite concave faces, as illustrated in Figure 5.15.

Label: A *perturbing agent* is identified by a label placed in an unbordered box containing a string of characters. The characters can be distributed on several lines to improve readability, although this is not mandatory. The label 25

box must be attached to the center of the *perturbing agent* container. The label may spill outside of the container.




Figure 5.15: The Process Description glyph for *perturbing agent*.

Changes from Previous Version

No changes from pervious version.

StatefulEPN

5

Stateful entity pools can undergo physical changes, for example chemical modification or conformational change, which we wish to record in a Process Description map. This information is captured via the *StateVariable* (as can be seen in figure 5.16). The *LabellecCloneMarker* must be used to indicated that the *StatefulEPN* is cloned.

10

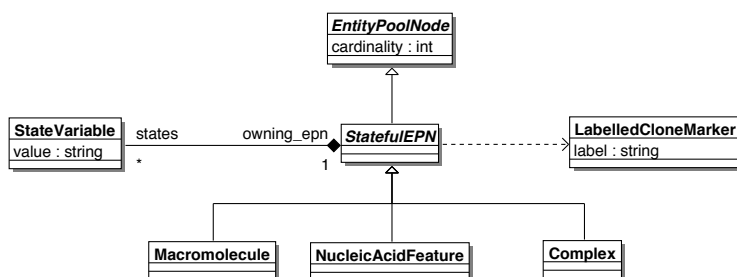


Figure 5.16: UML definition of the stateful entity pool node: showing its descendants and its association with state variables.

Generalisation

- EntityPoolNode (see section 5.5)

Attributes

No additional attributes.

Associations

15

states:StateVariable (*) The state variables that belong to this class.

Rules and Constraints

- State variables do not need to be logically unique, therefore two or more state variables with the same name are permitted.
- The `LabelledCloneMarker` must be used to indicate cloning for instances of `StatefulEPN` and its subclasses, with a must use the same

5

Changes from Previous Version

Not defined explicitly in the previous version.

Macromolecule

Many biological processes involve *macromolecules*: biochemical substances that are built up from the covalent linking of pseudo-identical units. Examples of macromolecules include proteins, nucleic acids (RNA, DNA), and polysaccharides (glycogen, cellulose, starch, etc.). Attempting to define a separate glyph for all of these different molecules would lead to an explosion of symbols in SBGN, so instead, SBGN Process Description Level 1 defines only one glyph for all macromolecules. The same glyph is to be used for a protein, a nucleic acid, a complex sugar, and so on. The exact nature of a particular macromolecule in a map is then clarified using its label and decorations, as will become clear below.

10

15

Generalisation

- `StatefulEPN` (see section 5.5)

Attributes

20

No additional attributes.

Associations

No additional associations.

Rules and Constraints

No additional rules and constraints.

25

Notation

There are two glyphs associated with `Macromolecule`. The first *Macromolecule monomer* is used when cardinality = 1 and the second *Macromolecule multimer* is used when cardinality > 1.


Glyph: *Macromolecule monomer***SBO Term:** SBO:0000245 ! macromolecule**Container:** A macromolecule is represented by a rectangular container with rounded corners, as illustrated in Figure 5.17.**Label:** A *macromolecule* is identified by a label placed in an unbordered box containing a string of characters. The characters can be distributed on several lines to improve readability, although this is not mandatory. The label box must be attached to the center of the container. The label may spill outside of the container. 5

Figure 5.17: The Process Description glyph for *macromolecule*, shown plain and unadorned on the left, and with an additional state variable and a unit of information in the middle and the cloned form on the right.

Glyph: *Macromolecule multimer*

10


SBO Term: SBO:0000420 ! multimer of macromolecules**Container:** A *multimer* is represented by two identical containers shifted horizontally and vertically and stacked one on top of the other. Figure 5.18 illustrates the glyph.**Label:** As monomer 15

Figure 5.18: The Process Description glyph for *macromolecule multimer*, shown plain and unadorned on the left, and with an additional state variable and a unit of information in the right and the cloned form on the right.

Usage Examples In this section, we provide examples of Entity Pool Node representations drawn using the SBGN Process Description Level 1 glyphs described above.

Figure 5.19 represents calcium/calmodulin kinase II, with phosphorylation on the sites threonine 286 and 306, as well as catalytic and autoinhibitory domains. 20
 Note the use of *units of information* and *state variables*.




Figure 5.19: An example representation of calcium/calmodulin kinase II.

Figure 5.20 represents the glutamate receptor in the open state, with both phosphorylation and glycosylation. The entity carries two functional domains, the ligand-binding domain and the ion pore, and its chemical nature is precided.




Figure 5.20: An example of a glutamate receptor in the open state.

Changes from Previous Version

No changes from the previous version.

5

NucleicAcidFeature

The `NucleicAcidFeature` represents a fragment of a macromolecule carrying genetic information. A common use for this construct is to represent a gene or transcript. The label of this EPN and its `natures` are often important for making the purpose clear to the reader of a map.

10

Generalisation

- `StatefulEPN` (see section 5.5)

Attributes

No additional attributes.

Associations

No additional associations.

15

Rules and Constraints

No additional rules and constraints.

Notation

The `NucleicAcidFeature` has two associated glyphs. The first *Nucleic acid feature monomer* is used when cardinality = 1 and the second, *Nucleic acid feature multimer* is used when cardinality > 1.

Glyph: *Nucleic acid feature monomer* This glyphs represents a monomeric macro- 5
molecule.

SBO Term: SBO:0000354 ! informational molecule segment

Container: A *nucleic acid feature* is represented by a rectangular container whose bottom half has rounded corners, as shown in Figure 5.21.

Label: The identity of a particular *Nucleic acid feature* is established by a label 10
placed in an unordered box containing a string of characters. The characters may be distributed on several lines to improve readability, although this is not mandatory. The label box must be attached to the center of the container. The label may spill outside of the container.




Figure 5.21: The Process Description glyph for *nucleic acid feature monomer*, shown plain and unadorned on the left and with an additional state variable and a unit of information in the middle and the cloned form on the right.

Glyph: *Nucleic acid feature multimer* This glyphs represents a multimeric macro- 15
molecule.

SBO Term: SBO:0000419 ! multimer of informational molecule segments

Container: A *Nucleic acid feature multimer* is represented by two identical containers shifted horizontally and vertically and stacked one on top of the other. Figure 5.22 illustrates the glyph.

Label: As monomer glyph.




Figure 5.22: The Process Description glyph for *nucleic acid feature multimer*, shown plain and unadorned on the left and with an additional state variable and a unit of information in the middle and the cloned form on the right.

Changes from Previous Version

No changes from the previous version.

Complex

A **Complex** represents a biochemical entity composed of other biochemical entities, whether macromolecules, simple chemicals, multimers, or other complexes (figure 5.23). The **Complex** can describe its composition by the set of **Subunits** it contains (see figure 5.5). This description is entirely optional and is there to assist the user with a visual shorthand about the composition of the complex. 5

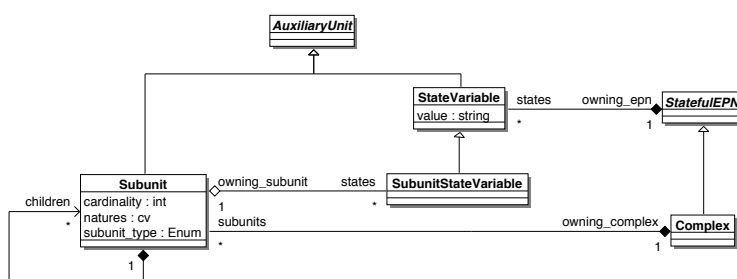


Figure 5.23: The UML definition of the **Complex** and its associated subunits. In particular this describes organisation of the state variables that belong to both the subunit, but also the complex.

Generalisations

- EntityPoolNode (see section 5.5)

10

Attributes

No additional attributes

Associations

subunits:Subunit (*) The subunits that describe the composition of this complex.

Special Rules and Constraints

15

- Once a set of subunits are defined for an **Complex** with a given **EntityType**, then they must be used by every instance using that entity type.¹⁴.
- The set of subunits in the **Complex** does not identify it. One or more **Complexes** that contain the same set of subunits, but have different labels are **not** identical.

20



See footnote.

¹⁴New rule.

Notation

The Complex is represented by two glyphs, the *Complex Monomer* which represents a Complex where the cardinality is one and the *Complex Multimer* where the cardinality is greater than that.

Complex Monomer

5

SBO Term: SBO:0000253 ! non-covalent complex

Container: A *complex* possesses its own container box surrounding the juxtaposed container boxes of its components. This container box is a rectangle with cut-corners (an octagonal box with sides of two different lengths). The size of the cut-corners are adjusted so that there is no overlap between the container and the components. The container boxes of the components must not overlap.

Label: The identification of a *named complex* is carried by an unbordered box containing a string of characters. The characters may be distributed on several lines to improve readability, although this is not mandatory. Ideally the label box should be attached to the midway between the border of the complex's container box and the border of the components' container boxes. However, if the Complex contains Subunit glyphs then the label may be positioned to optimise the clarity and avoid overlapping.




Figure 5.24: The *complex* glyph.

Complex Multimer

20

SBO Term: SBO:0000418 ! multimer of complexes

Container: A *Complex Multimer* is represented by two identical *Complex* containers shifted horizontally and vertically and stacked one on top of the other. Figure 5.25 illustrates the glyph.

Label: As monomer

25




Figure 5.25: The *Complex Multimer* glyph.

Layout Rules and Guidelines

- The subunits inside the complex must not overlap.
- The subunits should sit above the clone marker so that they are not obscured by it.
- The label should not be obscured by subunits or obscure them. 5

Changes from Previous Version

- Clarified that complex must have a label and the label identifies the complex irrespective of its subunit composition.
- The label positioning does not need to be at the centre of the Complex glyph.

Subunit 10



See footnote.

The Subunit¹⁵ is used to describe the composition of the Complex (see section 5.5). A complex can optionally be decorated with one or more subunits, which represent the types of EntityPoolNode (see section 5.5) that may aggregate to form a complex. As we can see from the UML representation (figure 5.23) the Subunit is an auxiliary unit that decorates the Complex and does not represent an entity pool directly. In addition it does not mimic the EntityPoolNode class hierarchy (Subunit, but rather uses the subunit_type attribute to indicate the type of subunit. 15

Generalisation

- EntityPoolNode (see section 5.5) 20

Attributes

cardinality: int (**R**) The number of copies of the subunit.

name: string (**O**) The name of the subunit.

subunit_type: enum (**R**) The type of the subunit. It can have one of the following values that correspond to the equivalent EPN class: SimpleChemical, UnspecifiedEntity, PerturbingAgent, Macromolecule, NucleicAcidFeature, Complex. 25

Associations

owning_complex:Complex (1) The complex that owns the subunit.

states:SubunitStateVariable (*) The state variables assigned to this subunit. 30

children:Subunit (*) Subunits that are contained by this subunit.

¹⁵The principles of the subunit have been agreed, but the details of this implementation should be reviewed. In particular the rules about state variables.

Rules and Constraints

- Two or more state variables with the same name are permitted.
- State variables with no name set are permitted.
- Subunits can also contain subunits. There is no limit on such nesting. The namespace rules below apply. 5
- The subunit defines a namespace for its state variables, e.g. subunit “A” assigned a state variable “P@Ser202” and a subunit “B” assigned the same state variable can be distinguished as A:P@Ser202 and B:P@Ser202.
- If the subunit is of type `Complex` then `children` can contain one or more `Subunit` instances. 10
- If the subunit has a cardinality > 1 then this should be displayed by the `AttributeValue` (see section 5.5).
- If `natures` contains one or more instances then these must be displayed via an `AttributeValue`.

Notation

15

The subunit symbol used for the *subunit* glyph varies depending on the `subunit_type` and `cardinality`. The symbols available are equivalent to those used by the EPN glyphs including the *complex*. Therefore it is possible to describe complexes within complexes. The mapping between these and the symbol used is shown table 5.1. Not that subunits may contain labels corresponding to their `name`. 20

Table 5.1: Mapping between the `subunit_type`, `cardinality` values of `Subunit` and the glyphs used to represent it. These are essentially the EPN glyphs described in this document.

subunit_type	cardinality = 0	cardinality > 0
SimpleChemical	<i>Simple Chemical Monomer</i>	<i>Simple Chemical Multimer</i>
UnspecifiedEntity	<i>Unspecified Entity</i>	None
PerturbingAgent	<i>Perturbing Agent</i>	None
Macromolecule	<i>Macromolecule Monomer</i>	<i>Macromolecule Multimer</i>
NucleicAcidFeature	<i>Nucleic Acid Feature Monomer</i>	<i>Nucleic Acid Feature Multimer</i>
Complex	<i>Complex Monomer</i>	<i>Complex Multimer</i>

The example in figure 5.26 illustrates the use of subunits in a complex. It also shows an equivalent complex without subunits. This is an import point. For every



See footnote.

Complex drawn with subunits it will always be possible to draw an equivalent version that does not use contains subunits¹⁶.




Figure 5.26: Both these complex glyphs are equivalent. The one on the left is described using subunit decorators, the one on the right describes the same thing without them.

Changes from Previous Version

In previous version of the spec the subunits of a **Complex** were regarded as an EPN. This however, is incorrect as it implies there are pools within pools, which breaks one of the fundamental paradigms of the Process Description language. This is corrected in the current version and subunits are now adornments of the **Complex**.

ProcessNode

The **Process** (figure 5.27) represents a process that transforms one or more entity pools into one or more entity pools, that are identical or different. A process may be used to represent or summarise more than one known process.

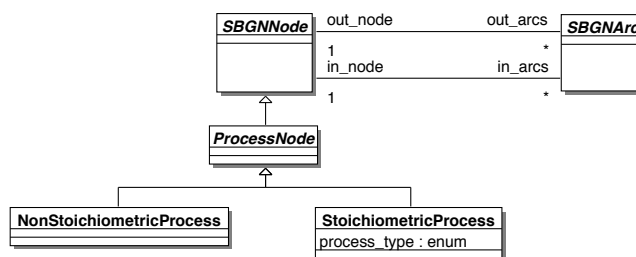


Figure 5.27: The UML definition of the **Process** and its associated subclasses. Note that the **Process** extends **SBGNNode** so all its descendants can potentially be nodes in a directed graph.

Generalisation

- **SBGNNode** (see section 5.5)

¹⁶Not discussed in detail. This must be true if states can be drawn on subunits, but actually belong to the complex. Either this or we enforce a rule that all state vars must be named uniquely.

Attributes

No additional attributes.

Associations

No additional associations.

Rules and Constraints

5

No additional rules and constraints.

Changes from Previous Version

- This was not explicitly defined in the previous version, but this version did define a glyph called *Process*. To avoid ambiguity this glyph has now been renamed *Stoichiometric Process* (see section 5.5). 10
- Previous specifications stated that processed could be duplicated when all associated EPNs were cloned. This behaviour has been changed the current status where all processes are unique in a Process Description map.

NonStoichiometricProcess



See footnote.

The NonStoichiometricProcess¹⁷ (figure 5.28) is a type of process. It does not necessarily result in a measurable change of entity pools, nor does it necessarily have a defined start and end point. In many cases the process is not well defined. This may be because it is not well understood or because the detail is not important or is being summarised. 15




Figure 5.28: The UML definition of the NonStoichiometricProcess and its associated subclasses.

Generalisation

20

- ProcessNode (see section 5.5)

Attributes

No additional attributes.

¹⁷This has been discussed and agreed in past meetings

Associations

No additional associations.

Rules and Constraints

No additional rules and constraints.

Changes from Previous Version

5

Not defined in the previous version.

Phenotype

A biochemical network can generate phenotypes or affect biological processes. Such processes can take place at different levels and are independent of the biochemical network itself. To represent these processes in a map, SBGN defines the Phenotype (figure 5.28). 10

Generalisation

- NonStoichiometricProcess (see section 5.5)

Attributes

name: string (**R**) The name of the phenotype. 15

Associations

No additional associations.

Logical Identity

Logical Key:

- owning_map 20
- name

Rules and Constraints

- The number of in_arcs must be > 0 .
- in_arc can only contain instances of ModulationArc (see section 5.5) and its subclasses. 25
- out_arcs must be empty.

Notation

Glyph: *Phenotype*

SBO Term: SBO:0000358 ! phenotype

Container: A *phenotype* is represented by an elongated hexagon, as illustrated in Figure 5.29.

Label: A *phenotype* is identified by a label placed in an unbordered box containing a string of characters. The characters can be distributed on several lines to improve readability, although this is not mandatory. The label box must be attached to the center of the *phenotype* container. The label may spill outside of the container.




Figure 5.29: The Process Description glyph for *phenotype*.

Changes from Previous Version

This definition clarifies that the Phenotype cannot be cloned as it is now a subclass of Process, which is always unique.

SubmapNode



See footnote.

The SubmapNode¹⁸ (figure 5.30) is a placeholder for another process and is used when one wishes to hide the detail of this process from the Process Description map, but make it available to the reader as a separate related map. The Submap is not equivalent to an OmittedProcess (section 5.5). The Submap allows the detail of section of the Process Description map to be exported to another Process Description map and replaced by the SubmapNode, which acts as a place-holder. This is described in section 5.5 and the semantics of submap linking is defined in section 5.8.




Figure 5.30: The UML definition of the SubmapNode and its relationship to its submap, tags etc.

¹⁸This name change has not been discussed at the time of writing. The aim is to provide clarity between the submap and this glyph.

Generalisation

- NonStoichiometricProcess (see section 5.5)

Attributes

name: string (**R**) The name of the submap that is being summarised. Note that this name ideally will indicate the function or the processes that are being summarised. 5

Associations

terminals:SubmapTerminal (1..*) The terminals provide a reference between the EPNs in the Main Map and those in the submap, which are identified by a Tag. 10

Logical Identity

Logical Key:

- owning_map
- name

Rules and Constraints

15

- All instances of SubmapTerminal (see section 5.5) held by this class must be logically unique.
- attribin_arcs and out_arcs must be empty (i.e., degree = 0).

Notation

Glyph: *Submap Node*

20

SBO Term: SBO:0000395 ! encapsulating process

Container: The *submap* is represented as a square box to remind the viewer that it is fundamentally a process.

Label: The identification of the *submap* is carried by an unbordered box containing a string of characters. The characters may be distributed on several lines to improve readability, although this is not mandatory. The label box has to be attached to the center of the container box. 25




Figure 5.31: The Process Description glyph for *submap*. (a) the basic glyph symbol, without the *submap terminal* auxiliary units that would normally be associated with it. (b) The glyph as it would typically be used within a map — associated with EPN glyphs and containing *submap terminals*.

Changes from Previous Version

This glyph was called *Submap* in previous version of the Process Description specification. This is confusing when talking about the Submap itself so this glyph is now referred to as the SubmapNode to distinguish it.

LogicalOperator

5




Figure 5.32: The UML definition of the LogicalOperator.

The LogicalOperator (figure 5.32) performs a Boolean operation on one or more inputs to give a binary output. The input must be a Boolean value, and are obtained from the LogicArc (see section 5.5) connected to the LogicalOperator. The output a two-value quantity, 0 for False and positive non-zero for True. This is required because the output of the LogicalOperator must be connected to either a LogicArc or a ModulationArc (see section 5.5) both of which require their out node to provide a quality. The behaviour of the logical operator for each type of operator_type is shown in the following table:

AND	All inputs must be True for output to be True, otherwise output is false.
OR	At least one input must be True for output to be True. If all inputs are False then output is False.
NOT	Only one input is permitted and the output is the inversion of the input. Therefore True gives False and False gives True.

Generalisation

15

- NonStoichiometricProcess (see section 5.5)

Attributes

operator_type: enum (**R**) The operator type must be one of the following enumerations: AND, OR, NOT.

Associations

No additional associations.

5

Rules and Constraints

- **in_arc** can only contain one or more instances of **LogicArc**.
- **out_arc** can only contain one or more instances of **LogicArc** or **Modulation-Arc**.
- if **operator_type** is AND or OR, then **in_arc** must contain two or more arcs. 10
- if **operator_type** is NOT then **in_arc** must contain only one arc.
- **out_arc** can contain only one arc.

Notation

Three glyphs are used to represent the different operator types. The glyphs are names after the corresponding type.

15

Glyph: *And*

SBO Term: SBO:0000173 ! and.

Node: *And* is represented by a circle carrying the word “AND”.




Figure 5.33: The Process Description glyph for *and*. Only two inputs are represented, but more would be allowed.

Glyph: *Or*

SBO Term: SBO:0000174 ! or.

20

Node: *Or* is represented by a circle carrying the word “OR”.




Figure 5.34: The Process Description glyph for *or*. Only two inputs are represented, but more would be allowed.

Glyph: *Not*

SBO Term: SBO:0000238 ! not.

Node: *Not* is represented by a circle carrying the word “NOT”.




Figure 5.35: The Process Description glyph for *not*.

Changes from Previous Version

Although the LogicOperator was not explicitly defined in the previous version the semantics and glyphs are unchanged. 5

StoichiometricProcess




Figure 5.36: The UML definition of the StoichiometricProcess. The class interacts with subclasses of FluxArc and ModulationArc.



See footnote.

A stoichiometric process¹⁹ produces a measurable change in the quantities of entity pools consumed and produced. This might imply modification of covalent bonds (conversion), modification of the relative position of constituents (conformational process) or movement from one compartment to another (translocation). Such a process will have a basal rate at which this change occurs, which can be affected positively or negatively by the other entity pools, which ‘modulate’ the process. Examples of this include stimulation, inhibition and catalysis. In an irreversible process the entity pools interacting with it can be grouped into inputs and outputs. However, a stoichiometric process can also be reversible and so for convenience we refer to these groupings as the “left-hand-side” (LHS) and “right-hand-side” (RHS) of the process²⁰ (figure 5.37). 10 15

¹⁹New concept, but discussed in previous meetings. The semantics of the process being stoichiometrically balanced has not been discussed in detail for the stoichiometric process, and this is the subject of a tracker query. The spec previously stated that the process should be balanced and this is therefore consistent with that.

²⁰Note this designation is purely for grouping and is used even then the sides of the reaction are above and below the process.




Figure 5.37: An illustration of the “sidedness” of a process. The designation of LHS and RHS is essentially arbitrary.

In the Process Description language this is represented by the **StoichiometricProcess** (figure 5.36). It can be one of several different types, which indicate the amount that is known about the process or in some cases the nature of the process, for example association and dissociation. The permitted values for `process_type` are described in the following table:

5

generic	A generic stoichiometric process that transforms a set of entity pools into another set of entity pools.
omitted	Omitted processes are processes that are known to exist, but are omitted from the map for the sake of clarity or parsimony. A single <i>omitted process</i> can represent any number of actual processes. The <i>omitted process</i> is different from a <i>submap</i> . While a <i>submap</i> references to an explicit content, that is hidden in the main map, the <i>omitted process</i> does not “hide” anything within the context of the map, and cannot be “unfolded”.
uncertain	Uncertain processes are processes that may not exist. A single <i>uncertain process</i> can represent any number of actual processes.
association	The association between one or more <i>EPNs</i> represents the non-covalent binding of the biological objects represented by those <i>EPNs</i> into a larger complex.
dissociation	The dissociation of an <i>EPN</i> into one or more <i>EPNs</i> represents the rupture of a non-covalent binding between the biological entities represented by those <i>EPNs</i> .

Since this process is stoichiometric the relative quantities of the entity pools participating the process must be specified. For this reason the **FluxArc** (see section 5.5) has an **stoichiometry** attribute and each **EntityPoolNode** (see section 5.5) has a cardinality, which should be balanced in a valid Process Description map. This is especially important where there is potential ambiguity in the stoichiometry of the process (figure 5.38).

10

A stoichiometric process is deemed to be reversible if its `in_arcs` are **FluxArcs** of type ‘reversible’ (see figure 5.39). Semantically, this permits a reversible flow of entities through the process. Modulation of a reversible process affects the rate of

15

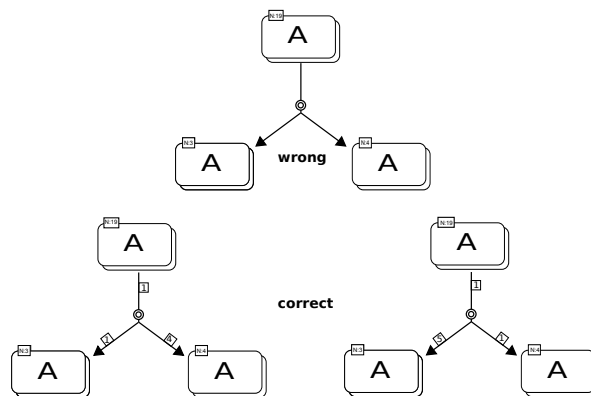


Figure 5.38: The figure illustrates why for the stoichiometry label is required to clarify potentially ambiguous stoichiometry. In the top example there is more than one possible solution, which can only be made clear using the stoichiometry labels in the bottom examples.

flux through the process, but does not directly affect the direction of that flow.




Figure 5.39: A valid reversible process. A process is reversible if its LHS and RHS contain only *production* arcs.

Generalisation

- ProcessNode (see section 5.5)

Attributes

process_type: enum (**R**) This must be one of the following enumerations: generic, 5
omitted, uncertain, association, dissociation.

Associations

No additional associations.

Rules and Constraints

General

10

- The `in_arc` must contain one or more `FluxArcs` containing the same `flux_type` value.
- The `in_arc` may only contains `FluxArc` instances with a `flux_type` of ‘consumption’, or ‘reversible’.
- In addition the `in_arc` may contain zero, one or more instances of `ModulationArc`. 5
- The `out_arc` must contain one or more instances of `FluxArc` with a `flux_type` or ‘production’.
- If `in_arcs` contains one or more `FluxArcs` of type ‘reversible’ this process reversible. 10
- The `EntityPoolNodes` that make up the LHS of the process should be consistent with the RHS, i.e. the process should be stoichiometrically balanced.²¹
- If at least one `FluxArc` associated with a `StoichiometricProcess` displays its stoichiometry via a *stoichiometry label* then all must.²²
- If more than one set of stoichiometries can be applied to the flux arcs of the process then the stoichiometry of the flux arcs must be displayed. 15



See footnote.



See footnote.

Association These rules apply if the `process_type` is ‘association’.

- The process must be irreversible.
- There can only be one ‘production’ `FluxArc`, with `stoichiometry = 1`.
- If a `Complex` is on the RHS of the association then there must be at least 2 EPNS on the LHS. *Is this too restrictive? It prevents multimers being represented as a complex of 2 identical subunits. It is taken from v1.0 of the spec and got lost in later versions.* 20

Dissociation These rules apply if the `process_type` is ‘dissociation’.

- The process must be irreversible. 25
- There can only be one ‘consumption’ `FluxArc`, with `stoichiometry = 1`.
- If a `Complex` is on the LHS of the dissociation then there must be at least 2 EPNS on the RHS. *see comment in association rules.*

²¹Tracker issue 329060. If the process is stoichiometric this must make sense. The previous spec states this so this is consistent with it.

²²Take from previous spec, but that said if one displays stoichiometry in a map which is too restrictive.

Notation

Glyph: *Process*

SBO Term: SBO:0000375 ! process

Node: A process is represented by a square box linked to two connectors: small arcs attached to the centers of opposite sides and referred to here as 'lugs'²³. The flux arcs are linked to the ends of the lugs as shown in figure 5.40. The lug's purpose is to 'gather' the flux arcs together before meeting the process node proper and in doing so they emphasize the 'sides' of the reaction. Therefore the lug must have a visually appreciable length²⁴ and must be placed on opposite sides of the process square. The modulatory arcs (section 5.5) point to the other two sides of the box.



See footnote.



See footnote.




Figure 5.40: The Process Description glyph for *process*.

The example in Figure 5.41 illustrates the use of a *process* node to represent the phosphorylation of a protein in a Process Description.




Figure 5.41: Phosphorylation of the protein MAP kinase.

The example in Figure 5.42 illustrates the use of a *process* node to represent a reaction between two reactants that generates three products.

²³The term lugs is used in discussion. We haven't discussed this in detail or at least come to a consensus on it. In particular does the lug need to be perpendicular to the process and does it need to be a straight line? How should it be used when the arc connecting to it is curved.

²⁴Undefined previously, but if we define it then it should be visible.




Figure 5.42: Reaction between ATP and fructose-6-phosphate to produce fructose-1,6-biphosphate, ADP and a proton.

The example in Figure 5.43 illustrates the use of a *process* node to represent a translocation. The large round-cornered rectangle represents a compartment border (see Section 5.5).




Figure 5.43: Translocation of calcium ion out of the endoplasmic reticulum. Note that the *process* does not have to be located on the boundary of the *compartment*. A *process* is not attached to any *compartment*.

The example in Figure 5.44 illustrates the use of a *process* node to represent the reversible opening and closing of an ionic channel in a Process Description.

5




Figure 5.44: Reversible opening and closing of an ionic channel.

When such a reversible process is asymmetrically modulated, it must be represented by two different processes in a Process Description. Figure 5.45 illustrates the use of two *process* nodes to represent the reversible activation of a G-protein coupled receptor. In the absence of any effector, an equilibrium exists between the inactive and active forms. The agonist stabilises the active form, while the inverse agonist stabilises the inactive form.

10




Figure 5.45: The reversible activation of a G-protein coupled receptor.

The example in Figure 5.46 presents the conversion of two galactoses into a lactose. Galactoses are represented by only one *simple chemical*, the cardinality being carried by the *consumption* arc.





Figure 5.46: Conversion of two galactoses into a lactose.

Glyph: *Omitted process*

SBO Term: SBO:0000397 - omitted process.

5


Node: An *omitted process* is represented by a *process* in which the square box contains a two parallel slanted lines oriented northwest-to-southeast and separated by an empty space.

Figure 5.47: The Process Description glyph for *omitted process*.**Glyph: *Uncertain process***

SBO Term: SBO:0000396 ! uncertain process.


10

Node: An *uncertain process* is represented by a *process* which square box contains a question mark.

Figure 5.48: The Process Description glyph for an *uncertain process*.

Glyph: Association**SBO Term:** SBO:0000177 ! non-covalent binding.

Node: An *association* between several entities is represented by a filled disc linked to two connectors, small arcs attached on point separated by 180 degrees. The consumption (Section 5.5) and production (Section 5.5) arcs are linked to the extremities of those connectors.

Figure 5.49: The Process Description glyph for *association*.

The example in Figure 5.50 illustrates the association of cyclin and CDC2 kinase into the Maturation Promoting Factor.




Figure 5.50: Association of cyclin and CDC2 kinase into the Maturation Promoting Factor.

Figure 5.51 gives an example illustrating the association of a pentameric macromolecule (a nicotinic acetylcholine receptor) with a simple chemical (the local anesthetic chlorpromazin) in an unnamed complex.




Figure 5.51: The association of a pentameric macromolecule with a simple chemical in an unnamed complex.

Glyph: Dissociation**SBO Term:** SBO:0000180 ! dissociation.

Node: A *dissociation* between several entities is represented by two concentric circles. A simple empty disc could be, in some cases, confused with the *catalysis* (section Section 5.5). Moreover, the existence of two circles reminds the dissociation, by contrast with the filled disc of the *association* (Section 5.5).




Figure 5.52: The Process Description glyph for *dissociation*.

The example in Figure 5.53 illustrates the dissociation of the small and large ribosomal subunits from a messenger RNA. 5




Figure 5.53: Dissociation of the small and large ribosomal subunits from a messenger RNA.

Changes from Previous Version

Although the `NonStoichiometricProcess` was not explicitly defined in the previous version the semantics and glyphs are unchanged.

Compartment

10




Figure 5.54: The UML definition of the `Compartment` showing how it contains `EntityPoolNode`.

The `Compartment` is a logical or physical structure that contains entity pool nodes. An `EntityPoolNode` (see section 5.5) can only belong to one compartment.

Therefore, the “same” biochemical species located in two different compartments are in fact two different pools.

Generalisation

- SBGNGlyph (see section 5.5)

Attributes

5

name: string (**R**) The name of the compartment.

natures: cv(*) A set of controlled vocabularies²⁵ that describes a characteristic of the compartment. Zero, one or more values may be set, but each one must belong to a different controlled vocabulary.



See footnote.

Associations

10

epns:EntityPoolNode (*) The EntityPoolNodes contained by this compartment.

Logical Identity

Logical Key:

- owning_map
- name

15

Rules and Constraints

- name must not be used by another instance of Compartment contained by the same instance of Map.
- epns must contain a unique set of EntityPoolNodes. See section 5.7 for the definition of EntityPoolNode uniqueness.

20

Notation

Glyph: *Compartment*

SBO Term: SBO:0000290 ! physical compartment

Container: A compartment is represented by a surface enclosed in a continuous border or located between continuous borders. These borders should be noticeably thicker than the borders of the EPNs. A compartment can take any geometry. A compartment must always be entirely enclosed. 25

Label: The identification of the compartment is carried by an unbordered box containing a string of characters. The characters can be distributed on several lines to improve readability, although this is not mandatory. The label 30

²⁵This reconciles the use of the Unit of Information to represent the nature of an EPN by using to present similar information for the compartment. This is consistent with previous usage, but not with the usage of the UofI for annotation.

box can be attached anywhere in the container box. Note that the label can spill-over from the container box.




Figure 5.55: The Process Description glyph for *compartment*.

To allow more aesthetically pleasing and understandable maps, compartments are allowed to overlap each other visually, but it must be kept in mind that this does not mean the top compartment contains part of the bottom compartment. 5
Figure 5.56 shows two semantically equivalent placement of compartments:




Figure 5.56: Overlapped compartments are permitted, but the overlap does not imply containment.

Overlapped (hidden) part of the compartment should not contain any object which could be covered by an overlapping compartment. Figure 5.57 illustrates the problem using an incorrect map.




Figure 5.57: Example of an **incorrect** map. Overlapped compartments must not obscure other objects.

Changes from Previous Version

The use Compartment has a set of natures, which previously were less well specified and handled as notation provided by the *unit of information*. In some cases, where the CVs used are not distinct or if the *unit of information* contains arbitrary text as annotation then maps containing these features will be invalid according to the current specification. 5

AttributeValue




Figure 5.58: The UML definition of the AttributeValue and its usage by other classes.



See footnote.

The AttributeValue²⁶ is used to present the values of certain attributes held by other SBGN elements. It is typically contained and owned by the class containing the attribute (or its descendants). It contains two values, one is a **code** to indicate the attribute that is defined and the other is the **value** itself. The **code** and the presentation format of the **value** are defined by the SBGN element that contains the AttributeValue, currently Compartment (see section 5.5), EntityPoolNode (see section 5.5), and Subunit (see section 5.5). 10

Generalisation

- AuxiliaryUnit (see section 5.5)

²⁶A new concept, that modifies the behaviour of the Unit of Information in previous versions with the need to use it to present the nature and cardinality of an EPN. The glyph retains its original name, but the class has been names to reflect it purpose.

Attributes

code: string (**R**) The code indicating the attribute that is being presented.

value: object (**R**) The value of the attribute. The format of the value is determined by the class holding the attribute.

Associations

5

No additional associations.

Rules and Constraints

No additional rules and constraints.

Notation

For historical reasons the `AttributeValue` is represented graphically by the glyph *Unit of Information*. 10

Glyph: *Unit of information* When representing biological entities, it is often necessary to convey some abstract information about the entity's function that cannot (or does not need to) be easily related to its structure. The *unit of information* is a decoration that can be used in this situation to add information to a glyph. Some example uses include: characterizing a logical part of an entity such as a functional domain (a binding domain, a catalytic site, a promoter, etc.), or the information encoded in the entity (an exon, an open reading frame, etc.). A *unit of information* can also convey information about the physical environment, or the specific type of biological entity it is decorating. 15 20

SBO Term: Not applicable.

Container: A unit of information is represented by a rectangle. The long side of the rectangle should be oriented parallel to the border of the *EPN* being annotated by the *unit of information*. The center of the bounding box of a *state of information* should be located on the mid-line of the border of the *EPN*. 25

Label: A *unit of information* is identified by a label placed in an unbordered box containing a string of characters. The characters can be distributed on several lines to improve readability, although this is not mandatory. The label box must be attached to the center of the container. The label may spill outside of the container. 30




Figure 5.59: The Process Description glyph for *unit of information*. (a) The glyph. (b) An example of its usage with a *macromolecule*.

Changes from Previous Version

There was no definition of the `AttributeValue` in the previous version of this specification. However, the *Unit of Information* did exist although its semantics have been changed. It no longer can hold arbitrary annotation but must display an attribute value and observe the constraints set out by the definition of the class owning the attribute. 5

Since the use of the *Unit of Information* has been deprecated, it is recommended that `Annotation` (see section 5.5) and the *Annotation* glyph is used instead.

StateVariable

10

Many biological entities such as molecules can exist in different *states*, meaning different physical or informational configurations. These states can arise for a variety of reasons. For example, macromolecules can be subject to post-synthesis modifications, wherein residues of the macromolecules (amino acids, nucleosides, or glucid residues) are modified through covalent linkage to other chemicals. Other examples of states are alternative conformations as in the closed/open/desensitized conformations of a transmembrane channel, and the active/inactive forms of an enzyme. 15

In the Process Description language these states are defined by the `StateVariableDefinitions` associated with the `EntityType`, but the specific values of the variables are define by the `StateVariable` (figure 5.60) associated with the `Entity-PoolNode`. For every `StateVariableDefinition` associated with an `EntityType` there should be a corresponding `StateVariable` associated with the instance of `Entity-PoolNode` using that type. This enforces one of the fundamental rules of the language that once a state variable has been displayed for a given entity type, then it must always be displayed. 20 25

Generalisation

- `AuxiliaryUnit` (see section 5.5)




Figure 5.60: The UML definition of the `StateVariable` showing its relationship to `StatefulEPN` and `StateVariableDefinition`.

Attributes

value: string (R) The value of the state variable. This is optional, but cannot be an empty string, should start with a non-space character and end with a non-space character. It should also not include the '@' character²⁷.



See footnote.

Associations

5

owning_epn:StatefulEPN (1) The stateful EPN that owns the state variable.
definition:StateVariableDefinition (1) The definition of this state variable²⁸.



See footnote.

Rules and Constraints

No additional rules and constraints.

Notation

10

Glyph: *State variable*

SBO Term: Not applicable.

Container: A *state variable* is represented by a “stadium” container, that is two hemicircles of same radius joined by parallel segments, as shown in Figure 5.61. The parallel segment axis should be tangent to the border of the glyph of the *EPN* being modified by the *state variable*. The center of the bounding box of a *state variable* should be located on the mid-line of the border of the *EPN*. In previous versions of this specification the *state variable* was represented by an ellipse. This symbols is now **deprecated** in favour of the stadium symbol described above. New Process Description maps should not use the ellipse symbol.

Label: An unbordered box containing a string indicating the contents of the `StateVariable`. The style of labeling of *State Variables* encouraged by SBGN Process Description Level 1 is to combine a prefix representing the value of the variable with a suffix representing the variable's name. Prefix and suffix should be separated by the symbol '@', X@Y thus meaning *value X AT variable Y*. If name is undefined then only the value should be displayed and

²⁷Not specified is previous spec, but seems essential if we are to avoid ambiguity.

²⁸New concept. It aims to ensure that all state variables are defined in all EPNS of the same type.

the '@' character omitted. If both the **name** and **value** are undefined then the label should be empty (i.e., an empty string). The label of a *state variable* should, if possible, be displayed within the boundary of the glyph. In earlier versions of the SBGN specification it was permitted to separate the **name** and **value** into two unlabelled boxes and display the **name** box outside the *state variable* glyph. This is now **deprecated** and new Process Description maps should not use this notation. 5




Figure 5.61: Examples of the Process Description glyph for *state variable*. Included are the older forms of glyph, which are now deprecated.

A *state variable* does not necessarily have to be Boolean-valued. For example, an ion channel can possess several conductance states; a receptor can be inactive, active and desensitized; and so on. As another example, a *state variable* “ubiquitin” could also carry numerical values corresponding to the number of ubiquitin molecules present in the tail. However, in all cases, a *state variable* on an EPN can only take *one* defined value. Further, an EPN’s *state variable* should always be displayed and always set to a value. An “empty” *state variable* is a *state variable* that is set to the value “unset”, it is not a *state variable* with no value. Note that the value “unset” is *not* synonymous to “any value” or “unknown value”. 10 15

Changes from Previous Version

The `StateVariable` class was not explicitly defined in previous versions of the specification, however the *state variable* was. Some aspects of its notation have been deprecated and these are detailed above (section 5.5). 20

Annotation

In SBGN Process Description Level 1 there are cases where the language does not capture everything the author wishes to convey. This may be additional experimental detail or descriptions of mechanisms that cannot be described full by the Process Description language. In this case the language provides the **Annotation**. This contains text and is associated with a particular glyph in a map. Importantly, it is purely “decoration” and does alter the meaning the map. 25

Generalisation

- AuxiliaryUnit (see section 5.5)

Attributes

annotation_text: string (**R**) The text of the annotation. The text is mandatory and cannot be empty or just spaces. 5

Associations

annotated_glyph:SBGNGlyph (1) The instance of SBGNGlyph that is being annotated²⁹.

Rules and Constraints

No additional rules and constraints. 10

Notation

Glyph: *Annotation*

SBO Term: SBO:NEW

Container: An *annotation* is represented by a rectangular container with a folded corner, as illustrated in Figure 5.62. This container is linked to the annotated element via a callout (see figure 5.63. The callout should overlap with the object it is annotating. 15

Label: An *annotation* contains information placed in an unbordered box containing a string of characters. The characters can be distributed on several lines to improve readability, although this is not mandatory. The label box must be attached to the center of the container. The label may spill outside of the container. 20




Figure 5.62: The Process Description glyph for *annotation*.

²⁹Note that as a result of this association only glyphs and **not** auxiliary items may be annotated by instances of Annotation




Figure 5.63: Example of *annotations* adding information to the description of the trans-phosphorylation of CaMKII. Note that three different types of links are used between annotation nodes and annotated elements. However, it is recommended to use a consistent scheme within a map.

Changes from Previous Version

This is a new language element and an not previous versions of the Process Description language.

CrossReference




Figure 5.64: The UML definition of the **Crossreference** showing its subclasses **Tag** and **SubmapTerminal** and its association with other elements in the Process Description language.



See footnote.

CrossReference³⁰ handles links or relationships between elements of a map and sub-map. At present there is only one reference glyph, *tag*, which can be used in a map referred to by a *submap* (Section 5.5) or as an auxiliary unit on the *submap*. The *clone marker* can also provide additional reference mechanisms and is discussed below (Section 5.5).

5

Generalisation

10

None

Attributes

reference_id: string (R) a string that identifies the cross-reference. The string cannot start and end in white space and cannot be empty.

³⁰A new concept required for consistency between tags and submap terminals.

Associations

equivalence:EquivalenceArc (1) The equivalence arc that links this class to the referenced element.

Rules and Constraints

- Two or more instances of `CrossReference` with the same `reference_id` value are pointing to the same element. 5
- The above rules applies within a Process Description map's namespace (see section 5.8).

Changes from Previous Version

Not defined in the previous version. 10

SubmapTerminal



See footnote.

A `SubmapTerminal`³¹ (figure 5.30) is a named reference that is part of a `SubmapNode` (see section 5.5). It provides the reference that is the link to a tag in the submap that the `SubmapNode` refers to.

Generalisation

15

- `AuxiliaryUnit` (see section 5.5)
- `CrossReference` (see section 5.5)

Attributes

No additional attributes.

Associations

20

No additional associations.

Rules and Constraints

No additional rules and constraints.

³¹Not properly defined in previous versions of the spec. The rules described here are consistent with accepted usage.

Notation

Glyph: *Submap Terminal*

SBO Term: Not applicable.

Container: A *tag* is represented by a rectangle fused to an empty arrowhead. The flat edge opposite the arrowhead should be aligned to the edge of the *Submap* glyph and the connecting should connect to the middle of this face (see figure 5.65). 5

Label: A *tag* is identified by a label placed in an unbordered box containing a string of characters. The characters can be distributed on several lines to improve readability, although this is not mandatory. The label box must be attached to the center of the container. The label may spill outside of the container. 10




Figure 5.65: The Process Description glyph for *Submap Terminal*. This shows the basic glyph and its correct usage within a *Submap* glyph.

Changes from Previous Version

Clarified that the tag does not link a *Compartment*, but only instances of *EntityPoolNode*. 15

Tag

A *Tag* is a named handle, or reference, to another *EntityPoolNode*. *Tags* are used to identify those elements in *submaps* (Section 5.5).

Generalisation

- SBGNGlyph (see section 5.5) 20
- CrossReference (see section 5.5)

Attributes

No additional attributes.

Associations

No additional associations.

Rules and Constraints

- All values of `reference_id` must be unique within an instance of `Map`.

Notation

5

Glyph: *Tag*

SBO Term: Not applicable.

Container: A *tag* is represented by a rectangle fused to an empty arrowhead, as illustrated in Figure 5.66. The symbol should be linked to one and only one edge (i.e., it should reference only one EPN or compartment).

10

Label: A *tag* is identified by a label placed in an unbordered box containing a string of characters. The characters can be distributed on several lines to improve readability, although this is not mandatory. The label box must be attached to the center of the container. The label may spill outside of the container.

15



Figure 5.66: The Process Description glyph for *tag*.

Changes from Previous Version

Clarified that the tag does not link a `Compartment`, but only instances of `Entity-PoolNode`.

FluxArc

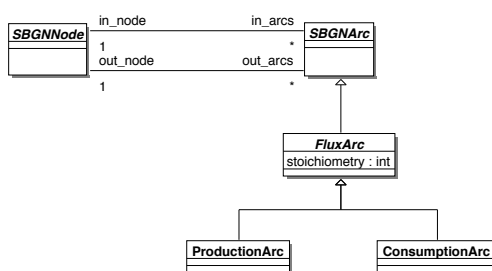


Figure 5.67: The UML definition of the `FluxArc` and its subclasses.

The **FluxArc** permits a quantity of entities to flow through the arc and in doing so connects a stoichiometric process (**StoichiometricProcess** (see section 5.5)) and an EPN (**EntityPoolNode** (see section 5.5)). The **FluxArc** has a stoichiometry which is used to indicate the stoichiometry of a process. It is required to eliminate ambiguity when the exact composition, or the number of copies, of the inputs or outputs to a reaction are ambiguous from the map (see figure 5.71 for an example). The **FluxArc** has three forms determined by its **flux_type**:

consumption Links an entity pool to a process that will consumed it (an input to the process).

production Links an entity pool to a process that will be produced by a process (an output from the process). It may also links entities on the “right-hand-side” of a reversible process (see section 5.5 for more details about the reversible process).

reversible Links a process to an entity pool that is the “left-hand-side” of a process (see section 5.5).

Generalisation

- **SBGNArc** (see section 5.5)

Attributes

stoichiometry: int (**R**) The stoichiometry of this **FluxArc** (see section 5.5). This must be a non-zero positive integer. **flux_type**: enumR The type of the flux arc. One of the following: consumption, production and reversible. No additional attributes.

Associations

No additional associations.

Rules and Constraints

- if the stoichiometry is not displayed then it is assumed to be 1³²
- If the stoichiometry > 1 then the stoichiometry must be displayed.³³
- if **flux_type** = ‘consumption’ or ‘reversible’ then:
 - The **in_node** must be an instance of **EntityPoolNode** (see section 5.5).



See footnote.



See footnote.

³²The previous version said that the stoichiometry was unspecified, but this was before the stoichiometric process was specified. In addition since cardinality is defined explicitly the stoichiometry can often be implied.

³³See above comment.

- The `out_node` must be an instance of `StoichiometricProcess` (see section 5.5).
- if `flux_type` = 'production' then:
 - The `in_node` must be an instance of `StoichiometricProcess` (see section 5.5).
 - The `out_node` must be an instance of `EntityPoolNode` (see section 5.5).

5

Notation

The `FluxArc` is represented by three glyphs depending on its `flux_type`:

consumption *consumption arc*.

10

production *production arc*.

reversible *reversible arc*³⁴.



See footnote.

These are defined below. In addition the *stoichiometry* is displayed by the *stoichiometry label*. Its appearance and layout in relation to the flux arc is also described here.

15

Glyph: *Consumption*

SBO Term: SBO:0000394 ! consumption.

End point: No particular symbol is used to represent a consumption.




Figure 5.68: The Process Description glyph for *consumption*.

Glyph: *Production*

SBO Term: SBO:0000393 ! production.

20

End point: The target extremity of a *production* carries a filled arrowhead.



Figure 5.69: The Process Description glyph for *production*.

³⁴The reversible arc is introduced here to fix syntax problems in previous spec. The change is visually neutral on previous Process Description maps.

Glyph: *Reversible*

SBO Term: SBO:0000393 ! production.

End point: The origin extremity of a *reversible* carries a filled arrowhead.




Figure 5.70: The Process Description glyph for *reversible*.

Glyph: *Stoichiometry Label* The stoichiometry label is part of the *consumption arc* and *production arc* glyphs see below (sections 5.5 and 5.5). However, as their use is common to all subclasses of FluxArc their presentation is described here. 5




Figure 5.71: Examples of stoichiometry label layout. In figure (a) the label is aligned with the stoichiometry box, while in (b) the label is aligned with the orientation of the map: these are both simple cases where the arc is a straight line. In cases where the arc is curved, the corners at the base of the label are anchored to point on the arc (c) and the label is drawn over the arc (d). Note that in (d) the covered part of the arc is shown for clarity, but normally the box is opaque and so the arc is not visible.



See footnote.

The label is a node that must be drawn above the flux arc. This node is attached to the arc where it intersects the arc with its bottom corners (see figure 5.71.)³⁵.

SBO Term: None

Container: A rectangle with a draw edge.

Label: A number that should remain within the container and be of a normal font, i.e., not bold or italic.

Changes from Previous Version



See footnote.

The *reversible arc* has been added to ensure the correct syntax is observed for a reversible process³⁶ — using the *production arc* for this was syntactically forbidden in the previous version of the specification. There is no change to semantics from 15

³⁵These layout rules were agreed at COMBINE 2010, but have not been voted on.

³⁶New glyph. We need it to ensure the correct syntax as we have a directed graph and we need to prevent the production arc pointing towards the process. It looks the same as the previous situation.

previous version, but layout rules for placement of stoichiometry label have been clarified.

ModulationArc




Figure 5.72: The UML definition of the **ModulationArc**. The class interacts with subclasses of **StoichiometricProcess**.

The **ModulationArc** (figure 5.72) affects the flux of a process represented by the target process. Such a modulation can affect the process **positively or negatively**, or even both ways depending on the conditions, for instance the concentration of the intervening participants. The permitted values for **process_type** are described in the following table:

modulation	A general modulation where the exact nature of the modulation is not specified or not known. Modulation can be used when one does not know the precise direction of the effect.
stimulation	A stimulation affects positively the flux of a process represented by the target process. This stimulation can be for instance a catalysis or a positive allosteric regulation. Note that <i>catalysis</i> exists independently in SBGN, see Section 5.5.
catalysis	A particular case of stimulation, where the effector affects positively the flux of a process represented by the target process. The positive effect on the process is due to the lowering of the activation energy of a reaction.
inhibition	An inhibition negatively affects the flux of a process represented by the target process. This inhibition can be for instance a competitive inhibition or an allosteric inhibition.
necessary_stim	A necessary stimulation, is one that is necessary for a process to take place. A process modulated by a necessary stimulation can only occur when this necessary stimulation is active.

As discussed in Chapter 2, it is implied, but not defined explicitly that the process has a rate at which it converts its LHS EPNs to its RHS EPNs (and vice-versa in the case of a reversible process). This concept is important in understanding how the Process Description language describes process modulation³⁷.



See footnote.

³⁷This is taken from the grammar section of the previous specification.

1. A *process* with no modulations has an underlying “basal rate” which describes the rate at which it converts inputs to outputs.
2. A *modulation* changes the basal rate in an unspecified fashion.
3. A *stimulation* is a modulation that increases the basal rate.
4. An *inhibition* is a modulation that decreases the basal rate. 5
5. The above types of modulation, when assigned to the same process, are combined and have a multiplicative effect on the basal rate of the process.
6. Modulators that do not interact with each other in the above manner, should be drawn as modulating different process nodes. Their effect is therefore additive. 10

Generalisation

- EntityPoolNode (see section 5.5)

Attributes

No additional attributes.

Associations

15

states:StateVariable (*) The state variables associated with this EPN.

Rules and Constraints

- At most one *necessary stimulation* can be assigned to a process node. Two *necessary stimulations* would imply an implicit AND or OR operator. For clarity only one *necessary stimulation* can be assigned to a *process*, and such combinations must be explicitly expressed using *logical operators*³⁸. 20
- At most one *catalysis* can be assigned to a *process*. Modulation by a catalysis arc implies that the exact biochemical mechanism underlying the process is known. In this context two *catalysis* cannot be assigned to the same process node as they represent independent reactions. Other EPNs can be assigned to the same process as a catalysis, such as modulators, stimulators, and inhibitors, and will have a multiplicative modulation on the reaction rate defined by the catalysis³⁹. 25



See footnote.



See footnote.

³⁸Taken from the grammar section of the previous specification.

³⁹Taken from the grammar section of the previous specification.

Notation

The ModulationArc is represented by a number of glyphs depending on its modulation_type. The table below defines what glyph is used for each type.

Type	Glyph
modulation	<i>Modulation</i>
stimulation	<i>Stimulation</i>
inhibition	<i>Inhibition</i>
necessary_stim	<i>Necessary Stimulation</i>

Glyph: *Modulation*

5

SBO Term: SBO:0000168 ! control.

End point: The target extremity of a *modulation* carries an empty diamond.



Figure 5.73: The Process Description glyph for *modulation*.

Figure 5.74 represents the effect of nicotine on the process between closed and open states of a nicotinic acetylcholine receptor. High concentrations of nicotine open the receptor while low concentrations can desensitize it without opening.

10







Figure 5.74: Modulation of nicotinic receptor opening by nicotine.

Glyph: *Stimulation***SBO Term:** SBO:0000170 ! stimulation.**End point:** The target extremity of a *stimulation* carries an empty arrowhead.Figure 5.75: The Process Description glyph for *stimulation*.**Glyph: *Catalysis*****SBO Term:** SBO:0000172 ! catalysis.**Node:** The target extremity of a *catalysis* carries an empty circle.Figure 5.76: The Process Description glyph for *catalysis*.**Glyph: *Inhibition*****SBO Term:** SBO:0000169 ! inhibition.**Node:** The target extremity of an *inhibition* carries a bar perpendicular to the arc.Figure 5.77: The Process Description glyph for *inhibition*.**Glyph: *Necessary stimulation*****SBO Term:** SBO:0000171 ! necessary stimulation.**Node:** The target extremity of a *necessary stimulation* carries an open arrow (to remind that it is a *stimulation*) coming after a larger vertical bar.Figure 5.78: The Process Description glyph for *Necessary Stimulation*.

Examples The example in Figure 5.79 below describes the transcription of a gene X, that is the creation of a messenger RNA X triggered by the gene X. The creation of the protein X is then triggered by the mRNA X. (Note that the same example could be represented using the gene as reactant and product, although it is semantically different.)




Figure 5.79: The creation of a messenger RNA X triggered by the gene X.

The example in Figure 5.80 below describes the transport of calcium ions out of the endoplasmic reticulum. Without IP3 receptor, there is not calcium flux, therefore, one cannot use a *stimulation*. The Necessary Stimulation instead represents this absolute stimulation.




Figure 5.80: The transport of calcium ions out of the endoplasmic reticulum into the cytosol. Note that IP3R crosses both compartment boundaries. This is allowed, but the Macromolecule should only belong to one of the compartments see section D.1 for more discussion of this issue.

Changes from Previous Version

5

The definition of `ModulationArc` did not exist in the previous version but there has been no changes to the glyphs and glyphs semantics in this version.

LogicArc

The `LogicArc` (figure 5.81) takes a quantity from either a `LogicalOperator` (see section 5.5) or an `EntityPoolNode` (see section 5.5) and converts it into a Boolean output, which serves as an input for a `LogicalOperator` (see section 5.5). How this

10

is done is not defined, but one could imagine that when a threshold value of the quantity is exceeded the output is True.




Figure 5.81: The UML definition of the LogicArc and its context.

Generalisation

- SBGNArc (see section 5.5)

Attributes

5

No additional attributes.

Associations

No additional associations.

Rules and Constraints

- The in_node must be an instance of EntityPoolNode or LogicalOperator. 10
- The out_node must be an instance of LogicalOperator.

Notation

Glyph: Logic arc Logic arc is used to represent the fact that an entity influences the outcome of a logic operator.

SBO Term: SBO:0000398 ! logical relationship.

15

End point: No particular symbol is used to represent a logic arc.




Figure 5.82: The Process Description glyph for logic arc.

Changes from Previous Version

No changes from the previous version.

EquivalenceArc

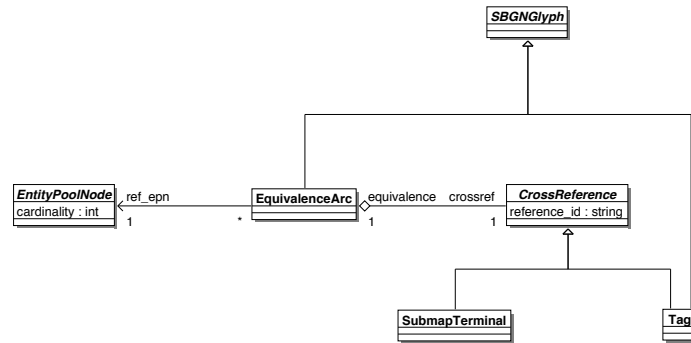


Figure 5.83: The UML definition of the EquivalenceArc and its context.

EquivalenceArc (figure 5.83) is the arc used to link a cross-reference to an EPN in another Process Description map (represented by CrossReference (see section 5.5)) with an EPN (EntityPoolNode (see section 5.5)) in this map.

Generalisation

5

- SBGNGlyph (see section 5.5)

Attributes

No additional attributes.

Associations

cross_ref:CrossReference (1) The cross reference associated to be associated with an EPN by this class. 10
 ref_epn:EntityPoolNode (1) The EPN that the cross-reference refers to.

Rules and Constraints

No additional rules and constraints.


Notation

15

Glyph: *Equivalence arc*

SBO Term: Not applicable.


End point: No particular symbol is used to represent an *equivalence arc*.

Figure 5.84: The Process Description glyph for *Equivalence arc*.

Changes from Previous Version

The relationship of *EquivalenceArc* to the *SubmapTerminal* (see section 5.5) was unclear in previous versions of the specification and has been clarified here.

CloneMarker

Figure 5.85: The UML definition of the *StateVariable* showing its relationship to *StatefulEPN*, *Complex* and *Subunit*.

If an *EntityPoolNode* (see section 5.5) is duplicated on a map, it is necessary to indicate this fact by the *CloneMarker* auxiliary unit (figure 5.85). The purpose of this marker is to provide the reader with a visual indication that this node has been cloned, and that at least one other occurrence of the *EntityPoolNode* can be found in the map (or in a submap; see Section 5.5). The clone marker takes two forms, simple and labeled, depending on whether the node being cloned can carry state variables (i.e., whether it is a stateful EPN). Note that an *EntityPoolNode* belongs to a single compartment. If two classes named “X” are located in two different compartments, such as ATP in cytosol and ATP in mitochondrial lumen, they represent different Entity Pools, and therefore do not need to be marked as cloned.

Generalisation

- *AuxiliaryUnit* (see section 5.5)

Attributes

No additional attributes.

Associations

owning_epn:EntityPoolNode (1) The EPN that holds this clone marker.

Rules and Constraints

No additional rules and constraints.

Changes from Previous Version

Not defined in previous version.

SimpleCloneMarker

5

The SimpleCloneMarker (figure 5.85) is the unlabelled subclass CloneMarker. All duplicated instances of StatelessEPN must contain an instance of this class.

Generalisation

- CloneMarker (see section 5.5)

Attributes

10

No additional attributes.

Associations

No additional associations.

Rules and Constraints

- Only subclasses of StatelessEPN (see section 5.5) can contain SimpleClone- 15
Markers.

Notation**Simple clone marker**

SBO Term: Not applicable.

Container: The simple (unlabeled) *clone marker* is a portion of the surface of an 20
EPN that has been modified visually through the use of a different shade,
texture, or color. Figure 5.86 illustrates this. The *clone marker* occupies the
lower part of the *EPN*. The filled area must be smaller than the unfilled one.

Label: Not applicable.




Figure 5.86: The Process Description glyph for *simple clone marker* applied to a *simple chemical* and a *multimer of simple chemicals*.

Figure 5.87 contains an example in which we illustrate the use of *clone markers* to clone the species ATP and ADP participating in different reactions. This example also demonstrates the chief drawbacks of using clones: it leads to a kind of dissociation of the overall network and multiplies the number of nodes required, requiring more work on the part of the reader to interpret the result. Sometimes these disadvantages are offset in larger maps by a reduction in the overall number of line crossings, but not always. In general, we advise that cloning should be used sparingly.




Figure 5.87: An example of using cloning, here for the species ATP and ADP.

Changes from Previous Version

No change from previous version.

LabeledClonerMarker

Unlike the `SimpleCloneMarker`, the `LabeledCloneMarker` (figure 5.85) includes (unsurprisingly, given its name) an identifying label that can be used to identify equivalent clones elsewhere in the map. This is particularly useful for subclasses of `StatefulEPN` (see section 5.5), because these can have a large number of

state variables displayed and therefore may be difficult to visually identify as being identical.

Generalisation

- CloneMarker (see section 5.5)

Attributes

5

label: string (**R**) The label that identified the clone. This label must start and end with an alphanumeric character, and cannot contain white space.

Associations

No additional associations.

Rules and Constraints

10

- At least two or more instanced of a **LabeledCloneMarker** with the same label must exist in this same in a given **Map** (see section 5.5).
- Only subclasses of **StatefulEPN** (see section 5.5) can contain labelled clone markers.

Notation

15

Labeled clone marker

SBO Term: Not applicable.

Container: The labeled *clone marker* is a portion of the surface of an *EPN* that has been modified visually through the use of a different shade, texture, or color. The *clone marker* occupies the lower part of the *EPN* glyph. The filled area must be smaller than the unfilled one, but the be large enough to have a height larger than the *clone marker's* label (cf below). 20

Label: A *clone marker* is identified by a label placed in an unbordered box containing a string of characters. The characters can be distributed on several lines to improve readability, although this is not mandatory. The label box must be attached to the center of the container. The label may spill outside of the container (the portion of the surface of the *EPN* that has been modified visually). The font color of the label and the color of the clone marker should contrast with one another. The label on a *labeled clone marker* is mandatory. 25 30




Figure 5.88: The Process Description glyph for *labeled clone marker* applied to a *macromolecule*, a *nucleic acid feature* and a *multimer of macromolecules*.

Changes from Previous Version

No changes from previous version.

5.6 Controlled vocabularies

Some classes in the SBGN Process Description language can contain particular kinds of textual annotation conveying information relevant to the class. Examples are the natures of an EntityPoolNode (see section 5.5) or Compartment (see section 5.5) or the value of the StateVariable (see section 5.5). The values held by these attributes can be taken from controlled vocabularies defined below. When displayed in some cases is mandatory to prefix a code indicating the type of controlled vocabulary used. This is in order to make it clear what the information is that the value refers to: for example 'mt' indicates that a value 'rna' is describing the material RNA.

In the rest of this section, we describe the controlled vocabularies (CVs) used in SBGN Process Description Level 1. They cover the following categories of information: an EPN's material type, an EPN's conceptual type, covalent modifications on macromolecules and the physical characteristics. These controlled vocabularies are *closed* in the sense that only the values defined for each CV can be used in a valid Process Description map and also closed because only the controlled vocabularies defined in this specification can be used in a valid Process Description map⁴⁰. We understand that this is of necessity restrictive, but in closing these definitions it means we can be clear about the meaning of all CV terms used in the specification. Updates to the CV terms and the CVs used are welcome and we encourage any changes or additions to be submitted as a tracker item at the address given on the front page of this specification.



See footnote.

Entity pool node material types

The material type of an EPN indicates its chemical structure. A list of common material types is shown in Table 5.2, but others are possible. The values are to be

⁴⁰The previous version was a bit vague but implied that the CVs were open. This has been changed after discussion on the list.

taken from the Systems Biology Ontology (<http://www.ebi.ac.uk/sbo/>), specifically from the branch having identifier SBO:0000240 (*material entity under physical entity representation*). The labels are defined by SBGN Process Description Level 1.

Name	Label	SBO term
Non-macromolecular ion	mt:ion	SBO:0000327
Non-macromolecular radical	mt:rad	SBO:0000328
Ribonucleic acid	mt:rna	SBO:0000250
Deoxribonucleic acid	mt:dna	SBO:0000251
Protein	mt:prot	SBO:0000297
Polysaccharide	mt:psac	SBO:0000249

Table 5.2: A sample of values from the *material types* controlled vocabulary (Section 5.6).

The material types are in contrast to the *conceptual types* (see below). The distinction is that material types are about physical composition, while conceptual types are about roles. For example, a strand of RNA is a physical artefact, but its use as messenger RNA is a role.

Entity pool node conceptual types

An EPN's *conceptual type* indicates its function within the context of a given Process Description. A list of common conceptual types is shown in Table 5.3, but others are possible. The values are to be taken from the Systems Biology Ontology (<http://www.ebi.ac.uk/sbo/>), specifically from the branch having identifier SBO:0000241 (*functional entity under physical entity representation*). The labels are defined by SBGN Process Description Level 1.

Name	Label	SBO term
Gene	ct:gene	SBO:0000243
Transcription start site	ct:tss	SBO:0000329
Gene coding region	ct:coding	SBO:0000335
Gene regulatory region	ct:grr	SBO:0000369
Messenger RNA	ct:mRNA	SBO:0000278

Table 5.3: A sample of values from the *conceptual types* vocabulary (Section 5.6).

Macromolecule covalent modifications

A common reason for the introduction of state variables (Section 5.5) on an entity is to allow access to the configuration of possible covalent modification sites on

that entity. For instance, a macromolecule may have one or more sites where a phosphate group may be attached; this change in the site's configuration (i.e., being either phosphorylated or not) may factor into whether, and how, the entity can participate in different processes. Being able to describe such modifications in a consistent fashion is the motivation for the existence of SBGN's covalent modifications controlled vocabulary. 5

Table 5.4 lists a number of common types of covalent modifications. The most common values are defined by the Systems Biology Ontology in the branch having identifier *SBO:0000210 (addition of a chemical group under interaction→process→biochemical or transport reaction→* 10

biochemical reaction→conversion). The labels shown in Table 5.4 are defined by SBGN Process Description Level 1; for all other kinds of modifications not listed here, the author of a Process Description must create a new label (and should also describe the meaning of the label in a legend or text accompanying the map).

Name	Label	SBO term
Acetylation	Ac	SBO:0000215
Glycosylation	G	SBO:0000217
Hydroxylation	OH	SBO:0000233
Methylation	Me	SBO:0000214
Myristoylation	My	SBO:0000219
Palmytoylation	Pa	SBO:0000218
Phosphorylation	P	SBO:0000216
Prenylation	Pr	SBO:0000221
Protonation	H	SBO:0000212
Sulfation	S	SBO:0000220
Ubiquitination	Ub	SBO:0000224

Table 5.4: A sample of values from the *covalent modifications* vocabulary (Section 5.6).

Physical characteristics 15

SBGN Process Description Level 1 defines a special unit of information for describing certain common physical characteristics. Table 5.5 lists the particular values defined by SBGN Process Description Level 1.

5.7 Entity Pool Node Identity and Cloning

All elements in an SBGN Process Description map have an implicit identify defined by the *id* attribute in *SBGNElement* (see section 5.5) so in that respect all elements drawn are unique. However, in some cases it is possible to draw two or 20

Name	Label	SBO term
Temperature	pc:T	SBO:0000147
Voltage	pc:V	SBO:0000259
pH	pc:pH	SBO:0000304

Table 5.5: A sample of values from the *physical characteristics* vocabulary (Section 5.6).

elements that define identical information in the map. In some classes, such as the `Compartment` this is prohibited, but in others (`EntityPoolNode`) it is not, but requires special decoration (the clone marker) to indicate that the information is replicated on the map. We define this type of identity and “logical identity” and the attributes that so discriminate between elements as the “logical key”⁴¹.

⚠
See footnote.

Clearly it is therefore important that we define the logical key for all the elements in the map and that we do this for the class definitions (section 5.5). The conventions used for this is as follows:

- If no logical key is specified then instance identity applies (see above).
- If the logical key is defined then this should be applied when determining equality between two elements.
- If an element can be replicated and two elements are determined to be logically identical, then instance identity should be used to determine uniqueness.
- If an element can be replicated if marked as a clone, then the above rule applies. If the element is not marked as cloned, then the only the logical key can be used when determining uniqueness.

Figure 5.89 illustrates how the above rules are applied in practise in the Process Description map. They also affect our understanding when reading a map. For example since all processes are unique, two *Process* glyphs connected to the same input and output indicates that the same “reaction” can be carried out by two distinct processes, with potentially different mechanisms.

5.8 Map and Submap Linking

The submap is a visual device that allows the detail of an Process Description map to be exported into another Process Description map and replaced by a `SubmapNode` (see section 5.5), which acts as a place-holder. In the case of an SBGN description that is made available through a software tool, the content of a `SubmapNode`

⁴¹This is recapitulation of the rules as defined in the previous version of the specification and as it is understood in discussions.




Figure 5.89: This fictitious example illustrates the correct cloning semantics in the Process Description language. The diagram shows the association of the phosphorylated macromolecule A (A[P]) and simple chemical B into a complex. Two copies of B are consumed and these are shown as separate subunits in the complex (note that no clone marker is used). The complex then dissociates into the same components that formed it which requires an appropriate clone marker to be applied. A[P] is translocated into compartment 2 where A[P] is not cloned as it does not have a duplicate. The unphosphorylated form of A dimerised, but again this does not require a clone marker as these are regarded as separate entity pools.

may be available to the tool. A user could then ask the tool to expand it, for instance by clicking on the *submap node* glyph. The tool might then expand and show its submap within the same map, or it might display it as a separate map. In the case of an SBGN description made available in a book or a website, the content of the *submap node* may be available on another page, possibly accessible via a hyperlink.⁴² 5



See footnote.

As you can see the rules relating to the submap concept range wider than those applying to the **SubmapNode** so we will detail these rules here. The concepts and nomenclature here can get confusing, so we will use the following nomenclature in the rest of this section⁴³: 10



See footnote.

supermap The map that owns the submap and so contains the submap glyph

⁴²Currently we can't name a map so linking between supermaps and submaps is problematic.

⁴³This nomenclature has not been used previously. The previous nomenclature was a bit confusing and didn't capture the hierarchy of maps well.

referring to it.

submap The map that contains the detail summarised by the submap glyph in the super-map.

reference_id the value displayed by the tag and submap terminal. These must be the same to indicate equivalence between map and submap. 5

Figure 5.90 represents a *submap node* that transforms glucose into fructose-6-phosphate. The *submap node* carries five terminals, four linked to EPNs and one linked to a *compartment*. The latter is particularly important in the case of EPNs present only in a *compartment* enclosed in a *submap node*, and that are not linked to terminals themselves. Note that the terminals do not define a “direction”, such as input or output. The flux of the reactions is determined by the context. 10




Figure 5.90: Example of a submap with contents elided.

The map in Figure 5.91 represents an unfolded version of a *submap*. Here, anything outside the *submap node* has disappeared, and the internal *tags* are not linked to the corresponding external *submap terminals*.⁴⁴



See footnote.

Namespace and the relationship between maps

15

The following rules apply to the use of the SubmapNode and submaps in the Process Description language⁴⁵.



See footnote.

- Submaps can contain SubmapNode instances so submaps can contain submaps - there is no limit on the depth of this.
- Submaps share the same namespace as their supermap. Therefore submaps of submaps share the top-most map's namespace. 20

⁴⁴This example in the previous version of the spec included compartment tagging. This has been removed, but the example may need refinement as a result.

⁴⁵These rules are taken from the previous version of the spec, but rules about the depth of submap 'recursion', the uniqueness of reference_id values and cloning between maps are not

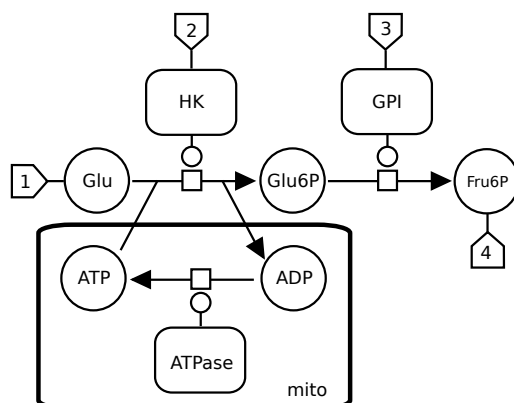


Figure 5.91: Example of an unfolded submap. The unfolded submap corresponds to the folded submap of Figure 5.90.

- Compartment's can be replicated across submaps and they all refer to the same compartment if they share the same name.
- Subclasses of `EntityPoolNode` replicated in both map and its immediate submap must be associated with a subclass of `CrossReference` (`SubmapTerminal` for the supermap and `Tag` for the submap). 5
- Subclasses of `EntityPoolNode` may be replicated between submaps separated by one or more other submaps, i.e. those not affected by the above rule.
- All `reference_ids` used must be unique within a given map namespace.
- submap folding: all tagged nodes in submap merge with those linked to terminals in super-map. If either node is cloned then all equivalent nodes in merged map must share the appropriate clone marker. 10
- For a valid mapping between subclasses of `EntityPoolNode` in a supermap and submap to exist, the `EntityPoolNode` subclasses must be logically identical and the `reference_id` in their associated `CrossReference` must be identical. identical. 15
- A `SubmapNode` which do not have a corresponding submap is valid.
- Submaps without a corresponding supermap are invalid⁴⁶.
- Subclasses of `EntityPoolNode` with the same name must share the same `EntityType`. 20



See footnote.

⁴⁶We need to sort out mapping between super and sub maps to get this working. I propose naming the maps

Compartment spanning

An *EPN* cannot *belong* to more than one *compartment*. However, an *EPN* can be *drawn* over more than one *compartment*. In such cases, the decision on which is the owning *compartment* is deferred to the drawing tool or the author. A *complex* may contain *EPNs* which belong to different *compartments* and in this way a *complex* can be used to describe entities that span more than one compartment. 5

This restriction makes it impossible to represent in a semantically correct way a macromolecule that spans more than one compartment — for example a receptor protein. It is clearly desirable to be able to show a macromolecule in a manner that the biologist expects (i.e. spanning from the outside through the membrane to the inside). Therefore, the author is recommended to draw the macromolecule across compartment boundaries, but the underlying SBGN semantic model will assign it to only one. The assignment to a *compartment* may be decided by the software drawing tool or the author. Note that this has implications for auto-layout algorithms as they will only be able to treat such *entity pool nodes* as contained within a *compartment* and will have no way of knowing a macromolecule spans a compartment. 10 15

The current solution is consistent with other Systems Biology representations such as SBML and BioPAX. For more information about the problems representing membrane spanning proteins and the rationale behind the current solution see Section D. 20

Layout Rules for a Process Description

6.1 Introduction

5

The previous chapters describe the appearance and meaning of SBGN Process Description Level 1 components. Here we describe rules governing the visual appearance and aesthetics of the Process Description language. The components of a Process Description have to be placed in a meaningful way – a random distribution with spaghetti-like connections will most likely hide the information encoded in the underlying model, whereas an elegant placement of the objects, giving a congenial appearance of the maps, may reveal new insights. The arrangement of components in a map is called a *layout*. 10

SBGN Process Descriptions should be easily recognisable not only by the glyphs used, but also by the general style of the layout. However, the arrangement of the components is a complex art in itself, and there is no simple rule which can be applied to all cases. Therefore this section provides rules for the layout of process description maps, divided into two categories: 15

1. requirements, i. e. rules which **must** be fulfilled by a layout, and
2. recommendations, i. e. rules which **should** be followed if possible. 20

In addition, we provide a list of additional suggestions which may help in producing aesthetically more pleasant layouts, possibly easier to understand.

Those layout rules are independent of the method used to produce the map, and apply to both manually drawn maps as well as maps produced by an automatic layout algorithm. The rules do not deal with interactive aspects (e. g. the effect of zooming). Further information about automatic network layout (graph drawing) can be found, for example, in the books of Di Battista and co-authors [?] and Kaufmann and Wagner [?]. 25

Please note that the color of objects do not carry any meaning in SBGN. Although one can use colors to emphasize part of a map or encode additional information, the meaning of the map should not depend on the colors. Furthermore, objects can have different sizes and size is also meaningless in SBGN. For example, a process node may be larger than a protein node. Also the meaning of a graph should be conserved upon scaling as far as possible. 5

6.2 Requirements

Requirements are rules which **must** be fulfilled by a layout to produce a valid Process Description map.

Node-node overlaps

10

Nodes are only allowed to overlap in two cases when they are allowed to contain other nodes — as described in Chapter 5. Otherwise, nodes are not allowed to overlap (Figure 6.1). This includes the touching of nodes. Touching is not allowed apart from the case where it has a specific meaning, e.g. two macromolecules touching each other within a complex because they form the complex. 15




Figure 6.1: Nodes must not overlap.

Node-edge crossing

Edges must be drawn on the top of a the node (Figure 6.2). See also recommendation Section 6.3.




Figure 6.2: If an edge crosses a node, the edge must be drawn on top of the node.

Node border-edge overlaps

Edges are not allowed to overlap the border lines of nodes (Figure 6.3). 20




Figure 6.3: Edges must not overlap node borders.

Edge-edge overlaps

Edges are not allowed to overlap (Figure 6.4). This includes touching of edges. Furthermore, an edge is neither allowed to cross itself nor to cross a boundary of node more than twice or other edges more than once.




Figure 6.4: Edges must not overlap.

Node orientation

5

Nodes have to be drawn horizontally or vertically, any other rotation of elements is not allowed (Figure 6.5).




Figure 6.5: The node orientation must be horizontally or vertically.

Node-edge connection

1. The arcs linking the square glyph of a *process* to the *consumption* and *production arcs* are attached to the center of opposite sides (Figure 6.6).
2. The modulatory arcs are attached to the other two sides, but not necessarily all to the center, as several modifiers can affect the same process node.

10




Figure 6.6: Arcs between a *process* and the *consumption* and *production arcs* must be attached to the center of opposite sides, modulatory arcs must be attached to the other two sides.

Node labels

At least a part of the label (unbordered box containing a string of characters) has to be placed inside the node it belongs to. Node labels are not allowed to overlap other nodes or other labels (this includes touching of other nodes or labels).

Edge labels

5

Edge labels are not allowed to overlap nodes. This includes touching of nodes.

Compartments

If a process has all participants in the same compartment the process node and all edges/arcs should be drawn in this compartment. If a process has participants in at least two different compartments, the process node has to be either in a compartment where the process has at least one participant or in the empty space.

10

6.3 Recommendations

Recommendations are rules which should be followed if possible and generally should improve the clarity of the diagram.

Node-edge crossing

15

Situations where edges and nodes cross should be avoided. Note that some crossings may be unavoidable, e. g. the crossing between an edge and a compartment border or an edge and a complex (if the edge connects an element inside the complex with something outside).

Labels

Labels should be horizontal. Node labels should be placed completely inside the node if possible. Edge labels should be placed close to the edge and avoid overlapping the edge as well as other edge labels.

Avoid edge crossings

5

The amount of crossings between edges should be minimized.

Branching of *association* and *dissociation*

The branching points of *association* and *dissociation* nodes should be placed closed to the symbol of the process, if possible at a distance comparable than, or smaller to, the diameter of the symbol defining the process (Figure 6.7).

10

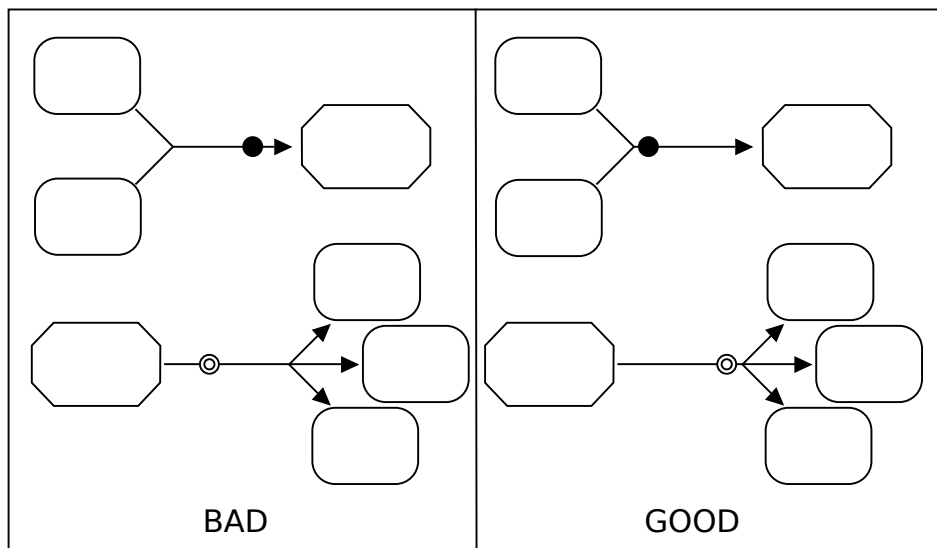


Figure 6.7: Branching points should be close to association and dissociation symbols.

Units of information

Units of information should not hide the structure of the corresponding node and should not overlap other elements (Figure 6.8).

6.4 Additional suggestions

Here is a list of additional layout suggestions which may help improve the aesthetics and clarity of Process Description maps. 15




Figure 6.8: Units of information should not overlap with any other element.

- Angle of edge crossings: If edge crossing cannot be avoided then the edges should cross with an angle close to 90 degrees.
- Drawing area and width/height ratio: The drawing should be compact and the ratio between the width and the height of the drawing should be close to 1. 5
- Edge length: Long edges should be avoided.
- Number of edge bends: Edges should be drawn with as few bends as possible.
- Similar and symmetric parts: Similar parts of a map should be drawn in a similar way, and symmetric parts should be drawn symmetrically. 10
- Proximity information: Related elements (e. g. nodes connected by a process or all elements within a compartment) should be drawn close together.
- Directional information: Subsequent processes (e. g. a sequence of reactions) should be drawn in one direction (e. g. from top to bottom or from left to right). 15
- Compartments: It can help clarity to use a different background shade or color for each compartment.

Acknowledgments

Here we acknowledge those people and organisations that assisted in the development of this and previous releases of the SBGN Process Description language specification. First we specifically acknowledge those who contributed directly to each revision of the specification document, followed by a comprehensive acknowledgement of contributors that attended workshops and forum meetings or in some other way provided input to the standard. Finally, we acknowledge the bodies that provided financial support for the development of the standard.

7.1 Level 1 Release 1.0

The specification of was written by Nicolas Le Novère, Stuart Moodie, Anatoly Sorokin, Michael Hucka, Falk Schreiber, Emek Demir, Huaiyu Mi, Yukiko Matsuoka, Katja Wegner and Hiroaki Kitano. In addition, the specification benefited much from the help of (in alphabetical order) Frank Bergmann, Sarala Dissanayake, Ralph Gauges, Peter Ghazal, Lu Li, and Steven Watterson.

7.2 Level 1 Release 1.1

The specification of SBGN PD Level 1.1 was written by Stuart Moodie and Nicolas Le Novère, with contributions from (in alphabetical order) Frank Bergmann, Sarah Boyd, Emek Demir, Sarala Wimalaratne, Yukiko Matsuoka, Huaiyu Mi, Falk Schreiber, Anatoly Sorokin, Alice Villéger.

7.3 Level 1 Release 1.2

The specification of SBGN PD Level 1.2 was modified by Stuart Moodie, with contributions from (in alphabetical order) Sarah Boyd, Nicolas Le Novère, Huaiyu Mi.

7.4 Level 1 Release 1.3

The specification of SBGN PD Level 1.3 was modified by Stuart Moodie, with contributions from (in alphabetical order), Tobias Czauderna, Nicolas Le Novère, Anatoly Sorokin.

7.5 Comprehensive list of acknowledgements

5

Here is a more comprehensive list of people who have been actively involved in SBGN development, either by their help designing the languages, their comments on the specification, help with development infrastructure or any other useful input. We intend this list to be complete. We are very sorry if we forgot someone, and would be grateful if you could notify us of any omission.

10

Mirit Aladjemm, Frank Bergmann, Sarah Boyd, Laurence Calzone, Melanie Courtot, Emek Demir, Ugur Dogrusoz, Tom Freeman, Akira Funahashi, Ralph Gauges, Peter Ghazal, Samik Ghosh, Igor Goryanin, Michael Hucka, Akiya Jouraku, Hideya Kawaji, Douglas Kell, Sohyoung Kim, Hiroaki Kitano, Kurt Kohn, Fedor Kolpakov, Nicolas Le Novère, Lu Li, Augustin Luna, Yukiko Matsuoka, Huaiyu Mi, Stuart Moodie, Sven Sahle, Chris Sander, Herbert Sauro, Esther Schmidt, Falk Schreiber, Jacky Snoep, Anatoly Sorokin, Jessica Stephens, Linda Taddeo, Steven Watterson, Alice Villéger, Katja Wegner, Sarala Wimalaratne, Guanming Wu.

15

The authors are also grateful to all the attendees of the SBGN meetings, as well as to the subscribers of the `sbgn-discuss@sbgn.org` mailing list.

20

7.6 Financial Support

The development of SBGN was mainly supported by a grant from the Japanese *New Energy and Industrial Technology Development Organization* (NEDO, <http://www.nedo.go.jp/>). The *Okinawa Institute of Science and Technology* (OIST, <http://www.oist.jp/>), the *AIST Computational Biology Research Center* (AIST 25 CBRC, <http://www.cbrc.jp/index.eng.html>) the British *Biotechnology and Biological Sciences Research Council* (BBSRC, <http://www.bbsrc.ac.uk/>) through a Japan Partnering Award, the European Media Laboratory (EML Research gGmbH, <http://www.eml-r.org/>), and the Beckman Institute at the California Institute of Technology (<http://bnmc.caltech.edu>) provided additional support for SBGN workshops. Some help was provided by the *Japan Science and Technology Agency* (JST, <http://www.jst.go.jp/>) and the *Genome Network Project* of the Japanese Ministry of Education, Sports, Culture, Science, and Technology (MEXT, <http://www.mext.go.jp/>) for the development of the gene regulation network aspect of SBGN, and from the *Engineering and Physi-* 30 *cal Sciences Research Council* (EPSRC, <http://www.epsrc.ac.uk/>) during the redaction of the specification. 35

Complete examples of Process Description Maps

The following maps present complete examples of SBGN Process Descriptions representing Biological processes. They by no mean exhaust the possibilities of SBGN Process Description Level 1. 5

Figure A.1 presents an example of metabolic pathway, that exemplifies the use of the *EPNs simple chemical, macromolecule, and clone marker*, the *PNs process*, and the *connecting arcs consumption, production and catalysis*. 10

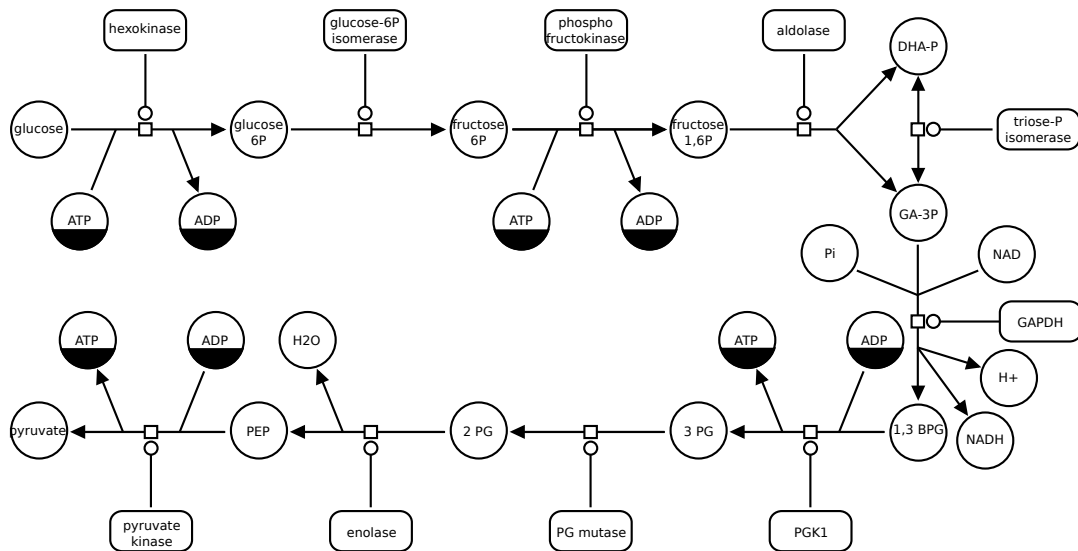


Figure A.1: Glycolysis. This example illustrates how SBGN can be used to describe metabolic pathways.

Figure A.2 presents an example of signalling pathway, that exemplifies in ad-

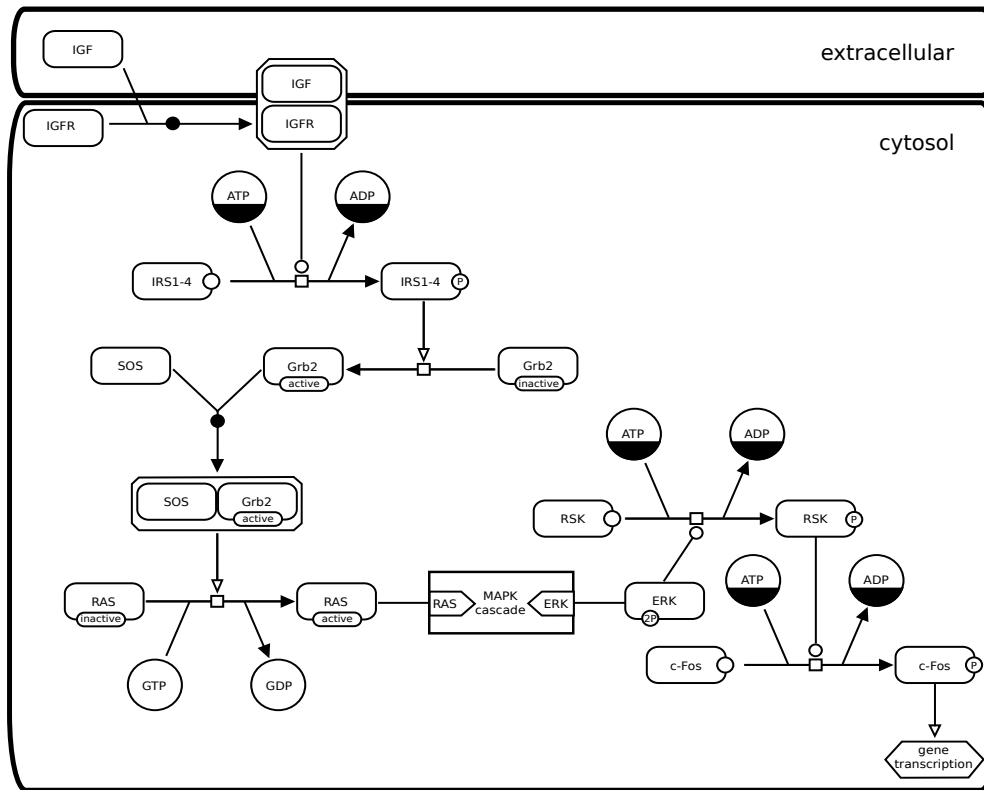


Figure A.2: Insulin-like Growth Factor (IGF) signalling. This example shows the use of compartments and how details can be hidden by using a submap. The submap is shown on Figure A.3.

dition the use of the *EPNs phenotype*, and *state variable*, the *containers complex*, *compartment* and *submap*, the *PNs association*, and the *connecting arcs stimulation*. Note the complex IGF and IGF receptor, located on the boundary of the compartment. This position is only for user convenience. The complex has to belong to a given compartment in SBGN Process Description Level 1.

Figure A.3 is an expanded version of the submap present on the map present in Figure A.2. It shows the use of *tag*.

Figure A.4 introduces an SBGN Process Description that spans several compartments. Note that the compartment “synaptic vesicle” is not **contained** in the compartment “synaptic button” but **overlaps** it. The *simple chemical* “ACh” of the “synaptic vesicle” is not the same *EPN* than the “ACh” of the “synaptic button” and of “synaptic cleft”. The situation is similar with the compartments “ER” and “muscle cytosol”. The map exemplifies the use of the *PN omitted* and *dissociation*, and the *connecting arc necessary stimulation*.

Figure A.5 introduces the use of SBGN Process Description Level 1 to encode

gene regulatory networks. It also show the use of the *EPNs Source* and the *logical operator and*.

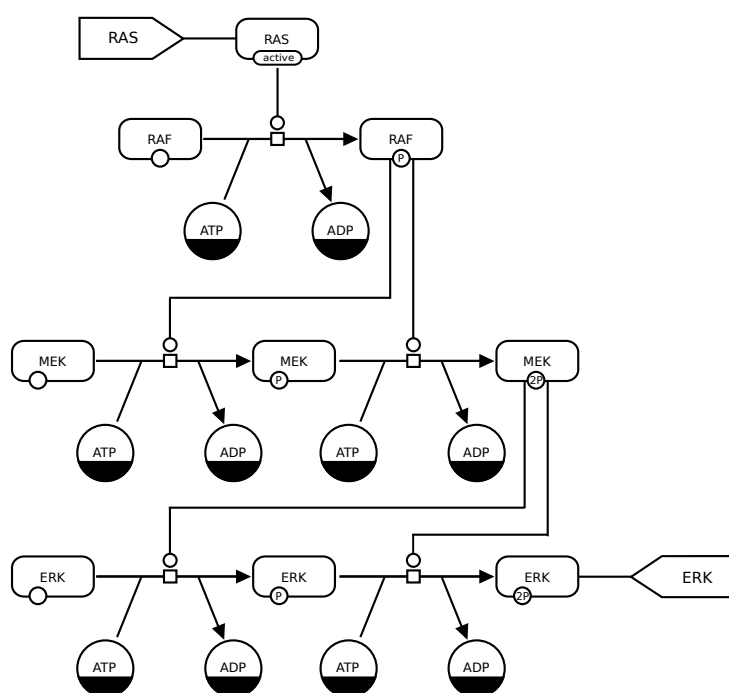


Figure A.3: A submap of the previous map showing the MAPK cascade.

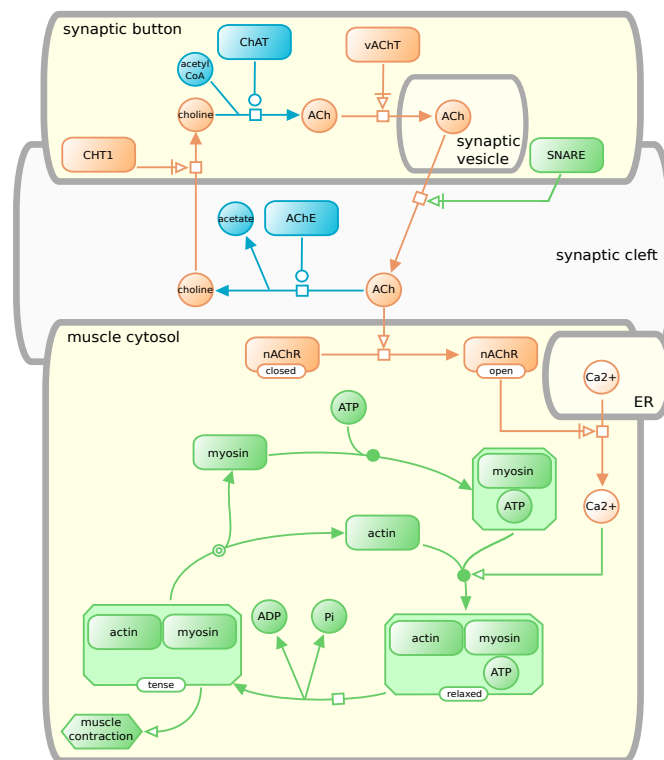


Figure A.4: Neuronal/Muscle signalling. A description of inter-cellular signalling using SBGN.

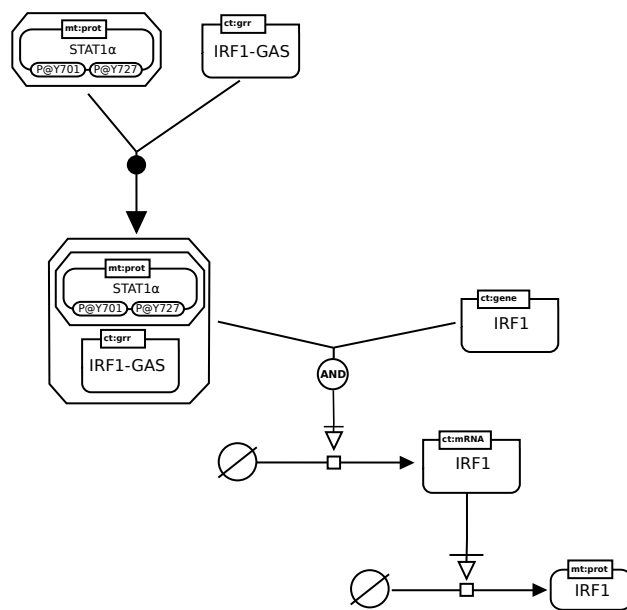



Figure A.5: Activated STAT1 α induction of the IRF1 gene. An example of gene regulation using logical operators.

Reference card

Print the summary of SBGN symbols on the next page for a quick reference.



Validation and consistency rules

This section describes criteria that must be met by the diagram (or SBGN-ML diagram description) to conform SBGN Process Description Level 1 specification. 5

All criteria has unique and perennial numbering. In the case when modification of the specification makes a rule irrelevant it will be deleted but the number of the rule will not be taken, so there will be gaps in rule numbering.

C.1 EPN rules 10

- 000001. Elements that do not comply SBGN Process Description Level 1 specification are not allowed on SBGN Process Description Level 1 map.
- 000002. All instances of **SBGNGlyph** (see section 5.5) must be unique (see section 5.7).
- 000003. The set of **SBGNNodes** linked to this node via a **SBGNArc** (its adjacent nodes) must be all belong to different entity pools (as defined by **Entity-PoolNode**) and cannot include more than one clone of the same entity pool. 15

Issues postponed to future levels

D.1 Multicompartment entities

5

The problem of entities, such as macromolecules, spanning several compartments proved to be a challenge for the community involved in the development of SBGN Process Description Level 1. It was thus decided to leave it for a future Level. It turns out there is at the moment no obvious solution satisfactory for everyone. Three broad classes of solutions have been identified so far:

10

- One can systematically locate an *EPN* in a given *compartment*, for instance a transmembrane receptor in a membrane. However, the reactions of this entity with entities represented by *EPN* in other compartments, such as extracellular ligands and second messenger systems, will create artificial transport reactions. 15
- One can represent the domains of proteins in different compartments by *macromolecules*, and link all those macromolecules in a *complex* spanning several compartments. However, such a representation would be very confusing, implying that the domains are actually different molecules linked through non-covalent bonds. 20
- One can accept *macromolecules* that span several compartments, and represent domains as *units of information*. Those *units of information* should then be located in given compartments. To make a full use of such a representation, one should then start and end connecting arcs on given *units of information*, something prohibited by the current specification. 25

D.2 Logical combination of state variable values

The value of a *state variable* has to be perfectly defined in SBGN Process Description Level 1. If a state variable can take the alternative values 'A', 'B' and 'C', one cannot attribute it values such as 'non-A', 'A or B', 'any' or 'none'. As a consequence some biochemical processes cannot be easily represented because of the very large number of state to enumerate. The decision to forbid such a Boolean logic lies in the necessity of maintaining truth path all over an SBGN map. 5

D.3 Non-chemical entity nodes

The current specification cannot represent combinations of events and entities. For instance a variable "voltage" cannot be controlled by a difference of concentration between different entities, such as a given ion in both sides of a membrane. 10

D.4 Generics

SBGN Process Description Level 1 does not provide mechanisms to sub-class *EPNs*. There is no specific means of specifying that *macromolecules* or *nucleic acid features* X1, X2 and X2 are subclasses of X. Therefore, any process that applies to all the subtypes of X has to be triplicated. That situation can easily generate combinatorial explosions of the number of *EPNs* or *PNs*. 15

D.5 State and transformation of compartments

In SBGN Process Description Level 1 a *compartment* is a stateless entity. It cannot carry *state variables*, and cannot be subjected to process modifying a state. As a result, a *compartment* cannot be transformed, moved, split or merged with another. If one want to represent the transformation of a compartment, one has to create the start and end compartments, and represent the transport of all the *EPNs* from one to the other. This is not satisfactory, and should be addressed in the future. 20
25

Revision History

E.1 Version 1.0 to Version 1.1

Below are the changes incorporated into Version 1.1 of the SBGN Process Description Level 1 specification. 5

Description	Tracker ID
Regarding modulation of reversible processes, changed “should” to “must” be represented by two <i>process</i>	
Removed “The connectors and the box move as a rigid entity” in the definition of <i>process</i>	
Changed the definition of process node to “represent processes that transform one or several EPNs into <i>one or several EPNs, identical or different</i> ”	
Changed SBO term of <i>compartment</i> From SBO:0000289 (functional compartment) to SBO:0000290 (physical compartment)	
Reorganised classification of glyphs	
Reorganised glyph section to reflect the above changes	
Revised reference card to reflect changes in glyph organisation	
Revised logic operators throughout spec to make sure input and output arcs meet before attaching to the glyph - as with processes.	
Added enumerated rules to grammar section. This is probably not complete, but should help the implementation of semantic validation by software tools. The hope is this will be refined as tools start validating maps.	
Updated UML maps and data dictionary to be consistent with rest of changes to spec.	
Definition of cardinality is ambiguous	2840996
<i>Sink and source</i> are lumped together	2726435
SBO terms are incorrect or missing.	2841261
continued on next page	

<i>continued from previous page</i>	
Description	Tracker ID
<i>Compartment</i> description is confusing and contradictory.	2841122
<i>Clone marker</i> fill percentages unhelpful.	2841114
Use of CV for physical characteristic not clear.	2841085
Definition of Cardinality is ambiguous.	2840996
input to AND on IFN example.	2804326
more SBO terms for <i>multimers</i>	2803593
Legend of figure 2.20 is incorrect	2803537
legend of figure 3.2	2802990
Compartment colouring	2745703
Errors in diag a4.	2664912
Change name of trigger glyph.	2664908
Transition should be renamed process.	2664862
Converting arcs tautological.	2664843
Example invalid.	2545870
consumption and production.	2388317
Should require circles to be distinguishable from ellipses	2219388
Figure 2.53	2162619
Reference card: production	2104471
Figure 2.42 is wrong	2104465
Mistake in the multi-cellular example	2395488
Should not prevent processes having identical in and out	2664933
No description of linking to subunit rules.	2545810
Extensively revised the grammar section. The UML diagrams have been simplified to show glyph taxonomy, and the data dictionary has been pruned to just show glyph identity. The some syntax rules have been moved into semantics and the rules reformulated to make them easier to understand.	
Eliminated duplicate rules in layout section and revised text slightly.	
Phenotype cloning?	2989007
Perturbing agent description	2940021

E.2 Version 1.1 to Version 1.2

Below are the changes incorporated into Version 1.2 of the SBGN Process Description Level 1 specification.

5	
Description	Tracker ID
Perturbing agent description	2940021
Members of complex touching	2849273
<i>continued on next page</i>	

<i>continued from previous page</i>	
Description	Tracker ID
PD Reference card error for submap glyph	3029242

E.3 Version 1.2 to Version 1.3

Below are the changes incorporated into Version 1.3 of the SBGN Process Description Level 1 specification.

Description	Tracker ID
Incorrect editor on title page	
Typos in acknowledgements	
Fixed typo in item on catalysis in section on modulation semantics.	
State variables figure 2.6 V1.2	3090543

Bibliography

- [1] M. Hucka, A. Finney, H. M. Sauro, H. Bolouri, J. C. Doyle, H. Kitano, A. P. Arkin, B. J. Bornstein, D. Bray, A. Cornish-Bowden, A. A. Cuellar, S. Dronov, E. D. Gilles, M. Ginkel, V. Gor, I. I. Goryanin, W. J. Hedley, T. C. Hodgman, J.-H. Hofmeyr, P. J. Hunter, N. S. Juty, J. L. Kasberger, A. Kremling, U. Kummer, N. Le Novère, L. M. Loew, D. Lucio, P. Mendes, E. Minch, E. D. Mjolsness, Y. Nakayama, M. R. Nelson, P. F. Nielsen, T. Sakurada, J. C. Schaff, B. E. Shapiro, T. S. Shimizu, H. D. Spence, J. Stelling, K. Takahashi, M. Tomita, J. Wagner, and J. Wang. The Systems Biology Markup Language (SBML): A medium for representation and exchange of biochemical network models. *Bioinformatics*, 19(4):524–531, 2003. 2825
- [2] Object management Group. OMG Unified Modeling Language (OMG UML), Superstructure, V2.1.2. Available via the World Wide Web at <http://www.omg.org/spec/UML/2.1.2/Superstructure/PDF/>, 2010. 2835
- [3] Michael Hucka, Frank T. Bergmann, Stefan Hoops, Sarah M. Keating, Sven Sahle, James C. Schaff, Lucian P. Smith, and Darren J. Wilkinson. The Systems Biology Markup Language (SBML): Language Specification for Level 3 Version 1 Core. Available via the World Wide Web at <http://sbml.org/Documents/Specifications>, November 2007. 2840
- [4] G. Di Battista, P. Eades, R. Tamassia, and I.G. Tollis. *Graph Drawing: Algorithms for the Visualization of Graphs*. Prentice Hall, New Jersey, 1998.
- [5] M. Kaufmann and D. Wagner. *Drawing Graphs: Methods and Models*, volume 2025 of *Lecture Notes in Computer Science Tutorial*. Springer, 2001.